

Merkle tree

Gianluca Dini

Dept. of Ingegneria dell'Informazione

University of Pisa

Email: gianluca.dini@unipi.it

Version: 07/04/2025

1

Brief history



UNIVERSITÀ DI PISA

- Ralph Merkle **patented Merkle Trees** in **1979**
- Merkle published the paper in 1987
 - R.Merkle. [A digital signature based on a conventional encryption function](#). CRYPTO 1987.
- Patent expired in 2002


apr. '25

Merkle Tree

2

2

Example: executable integrity [→]



UNIVERSITÀ DI PISA

- Executable X stored on disk as a list of blocks $\{x_1, x_2, \dots, x_r\}$
- OS needs to verify X integrity before execution


apr. '25

Merkle Tree

3

3

Example: executable integrity [→]



UNIVERSITÀ DI PISA

- Option 1
 - OS stores $t = H(X)$ on read-only memory^(*)
 - OS checks whether $t == H(X)$ before exection
 - Drawback: hashing the whole executable may slow down executable launching
 - (*) Read-only memory implementation
 - Separate system that provides data to requestors
 - Digital signature with private key offline *either by digitally sign hash or by HW trick that protects certain memory area.*


apr. '25

Merkle Tree

4

4

Example: executable integrity [↓]




UNIVERSITÀ DI PISA

- Option 2
 - For each block x_i , OS stores $t_i = H(x_i)$ on read-only memory
 - OS checks whether $t_i == H(x_i)$ when execution moves to x_i
 - Drawback: storage overhead to store t_i 's
- Can we do any better?

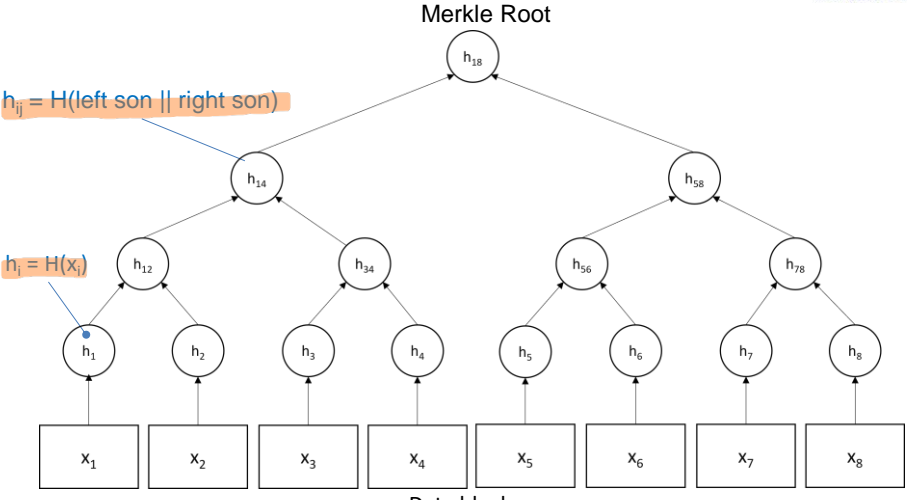
apr. '25 Merkle Tree 5

5

Merkle Tree (1979)



UNIVERSITÀ DI PISA



$h_{ij} = H(\text{left son} || \text{right son})$

$h_i = H(x_i)$


Merkle Root

Data blocks

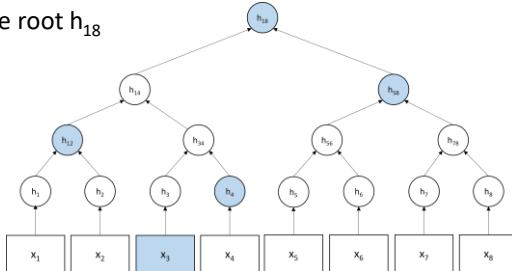
apr. '25 Merkle Tree 6

6 Assumption: tree is complete, otherwise you add leaves

Proving membership


 UNIVERSITÀ DI PISA

- Verify whether x_3 (contents and position) belongs to the data set
 - List of hashes: $\pi_3 = \langle h_4, h_{12}, h_{58}, h_{18} \rangle$ (Merkle proof)
 - Verification algorithm
 - Check whether $H(H(h_{12}, H(H(x_3), h_4)), h_{58}) == h_{18}$
 - Verify authenticity of the root h_{18}
- Complexity $O(\log n)$, with $n = \# \text{blocks}$
- The tree contains $2n - 1$ nodes (hashes)




Merkle Tree

apr. '25 7

7

• It is $\log n$ in storage and time. Block needs to be stored with merkle proof.

Properties



 UNIVERSITÀ DI PISA

- MT (or hash tree) allows efficient and secure verification of the contents of large data structures
- The root must be trusted: only thing I have to protect
 - Digitally signed
 - Maintained on a trusted source/storage
- Verifying whether a leaf node is part of the MT requires computing a #hashes proportional to the logarithm of the #leaves
 - $O(\log n)$, with n the number of leaves (blocks)
- MT does not require online secrets

apr. '25 8

8

Proving multiple membership



UNIVERSITÀ DI PISA

- Proving membership of multiple (L) elements
 - Trivial solution: L proofs $\rightarrow L \times \log_2 n$ hashes
 - Intuition: Many proofs overlap and thus can be shrinked


apr. '25

Merkle Tree

9

9

Proving multiple membership




UNIVERSITÀ DI PISA

- The proof
 - T set of blocks; $L \subseteq T$; $|L| = r$
 - $S = T - L$
 - $W = \text{cover}(S)$
 - Merkle proof = hash values corresponding to W
 - $W = \text{cover}(S)$ if every leaf in S is descendant of some node in W, and leaves outside S are not
- At most $r \cdot \log_2(n/r)$ hashes

apr. '25

Merkle Tree



T = B3, B7, B8
Merkle proof: h_{12}, h_4, h_{56}


Merkle Proof can be easily determined:
orange nodes are nodes that cover

I need specific blocks as proofs.

10

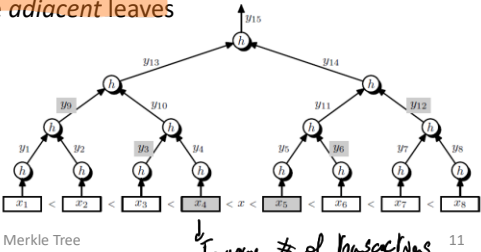
The white blocks without covering blue blocks.

Proving non-membership



UNIVERSITÀ DI PISA

- Proving non-membership
 - Sorted Merkle Tree: $x_1 < x_2 < \dots < x_n$ (e.g., set of TX's)
 - Prover provides
 - Proof π_i and proof π_j (merkle proof of adjacent nodes)
 - Verifier
 - Verifies π_i and π_j
 - Verifies that x_i and x_j are adjacent leaves
 - Verifies that $x_i < x < x_j$



Merkle Tree


11

apr. '25

11

The values must be sorted in one order, and I have to find adjacent values that contain x_i and verify if they are part of the tree and one adjacent.

Merkle Tree - applications



UNIVERSITÀ DI PISA

- File systems
 - IPFS, Btrfs, ZFS
- Content distribution protocols
 - Dat, Apache Wave
- Distributed revision control system
 - Git, Mercurial
- Blockchain
 - Bitcoin, Ethereum
- Backup Systems
 - Zeronet
- P2P networks
 - Torrent
- NoSQL systems
 - Apache Cassandra, Riak, Dynamo
- Certificate Transparency framework

apr. '25

Merkle Tree

13

13

Content distribution [→]

The diagram illustrates the content distribution process. A 'Content provider' box contains a vertical list of items [1, 2, ..., B] and a hash function 'H'. An arrow points from the list to 'H', which outputs a hash 'h_f'. A handwritten note 'store the hash on a trusted server' points to 'h_f'. A 'Trusted server' box contains a 'TS' block and also outputs 'h_f'. Below, 'Untrusted peers' are shown as circles p1, p2, ..., pn. One peer p2 is connected to a block blk_i, which is then connected to a user 'u'. A handwritten note 'download hash from trusted server and compare with your hash.' points to the user 'u'. The University of Pisa logo is in the top right.

apr. '25

Merkle Tree

14

14

Content Distribution [→]

The diagram shows the content distribution process. A 'Content provider' box contains a vertical list of items [1, 2, ..., B] and a hash function 'H'. An arrow points from the list to 'H', which outputs a hash 'h_f'. A handwritten note 'store the hash on a trusted server' points to 'h_f'. A 'Trusted server' box contains a 'TS' block and also outputs 'h_f'. Below, 'Untrusted peers' are shown as circles p1, p2, ..., pn. One peer p2 is connected to a block blk_i, which is then connected to a user 'u'. A handwritten note 'download hash from trusted server and compare with your hash.' points to the user 'u'. The University of Pisa logo is in the top right.

- How does the user know that the information that (s)he is getting from some untrusted peer is genuine and hasn't been tampered with (or corrupted)?


apr. '25

Merkle Tree

15

15

Content Distribution [→]



UNIVERSITÀ DI PISA

- **Solution no. 1 (shown in the slide)**
 - Trusted Server stores h_f
- **Verification**
 - Upon receiving all blocks $\{x_i, 1 \leq i \leq B\}$, compute $h_f' = H(x_1 \parallel x_2 \parallel \dots \parallel x_n)$.
 - Return $(h_f' == h_f)$
- **Drawback**
 - Check upon completion (possibly long delay)
 - Not possible to determine corrupted/compromised blocks


apr. '25

Merkle Tree

16

16

Content Distribution [→]



UNIVERSITÀ DI PISA

- **Solution n.2**
 - Trusted Server stores $\langle h_f, h_1, h_2, \dots, h_B \rangle$ with $h_i = H(x_i), 1 \leq i \leq B$
 - Number of hashes $B = \text{sizeof}(\text{file}) / \text{sizeof}(\text{block})$
 - Torrent: block size is 16 kbytes
- **User Verification**
 - The user can verify each block
- **Drawback**
 - Increase storage/bandwidth overhead


apr. '25

Merkle Tree

17

17

Content Distribution [↓]



UNIVERSITÀ DI PISA

- **Solution n.3: Merkle Tree**
 - **Trusted Server stores the root of the Merkle Tree**
 - **Each peer stores**
 - A **subset of the blocks $\{x_i\}$** ;
 - For **each block x_i** , a peer stores the corresponding Merkle proof π_i
 - **User Verification**
 - Upon **downloading a block x_i** , the user **verifies it using π_i and the tree root**

apr. '25


Merkle Tree

18

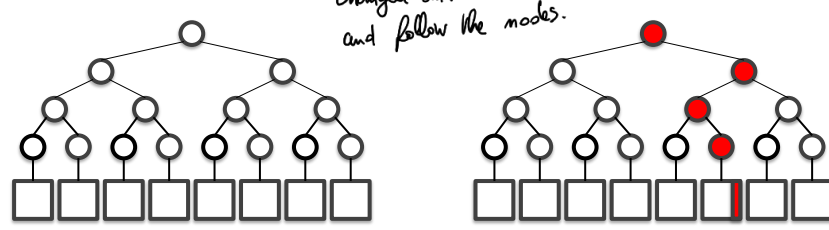
18

File comparison :

you can check whether file has changed and where: you compare root and follow the nodes.



UNIVERSITÀ DI PISA



- File F gets modified in a block x_i
- Comparing files takes is $O(B)$
- Comparing MTs is $O(\log B)$, **REPEAT VB?**

apr. '25

Merkle Tree

19

19

Replication

PC: Primary copy
SC: Secondary copy

Useful in replication systems in which you have primary and secondary copy of data. Suppose that one SC is disconnected and is not a perfect copy anymore, you can use which block has to be modified.

apr. '25

Merkle Tree

20

20

Replication

UNIVERSITÀ DI PISA

- How can the primary replica determine whether a disconnected secondary replica has to be updated?
- Upon reconnection, the primary replica compares its MT with the secondary replica's MT in order to determine the modified blocks

apr. '25

Merkle Tree

21

21