

# Defences

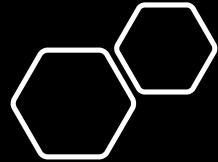
# What have we learned

Certain features of programming languages can introduce security vulnerabilities, making systems susceptible to attacks.

Today, we will explore possible countermeasures to mitigate these risks



A first defence: awareness



## Defensive Programming

We look out for trouble before it happens, anticipate the unexpected, and never put ourselves into a position from which we can't extricate ourselves

*Design by Contract:*  
clients and suppliers  
must agree on rights and responsibilities.

State precisely what your program does

# The Pragmatic Programmer

---

From Journeyman to Master

Andrew Hunt  
David Thomas

→ SWI, Kernels, Libraries, Languages, Runtimes all contribute to security of the ecosystem

## Security by construction

1. **Formal Methods & Verification** – Uses mathematical models to prove that a system adheres to security properties, reducing the likelihood of vulnerabilities.
2. **Secure Coding Practices** – Enforces strict adherence to secure programming techniques, such as memory safety and avoiding common pitfalls like buffer overflows.
3. **Least Privilege & Strong Typing** – Ensures that components and processes operate with minimal permissions, and strong type systems help prevent unintended behaviors.
4. **Design-Time Security Considerations** – Security is incorporated during system architecture and design phases rather than being bolted on later. (No patch on bad project)
5. **Automated Code Analysis & Tooling** – Uses static analysis, fuzz testing, and automated code reviews to detect potential security flaws early.
6. **Continuous Security Validation** – Implements secure development pipelines with ongoing security assessments.

↓  
Provide feedback asap

## DEFENCES (classical)

Apart from awareness, what else? Putting patches: if a feature you have in a PL is bad, you need this.

**Changing compilers and run-time mechanisms** to prevent exploitation of security vulnerabilities.

→ You will need to modify them

**This has extra work on compilers**

## DEFENCES (Innovative)

The design of evolutionary toolkits to support the design and development of secure programs enhances security within the software ecosystem by promoting design by construction.

This requires additional efforts in educating professionals to adopt secure design and development practices.

## DEFENCES (Innovative)

STOP using wrong prog. languages insecure, start with more secure ones.

**Design and implement novel programming languages  
where security and memory safety IS a main concern  
from the very beginning**

# Classical defences

# Mitigation 1: Security canaries

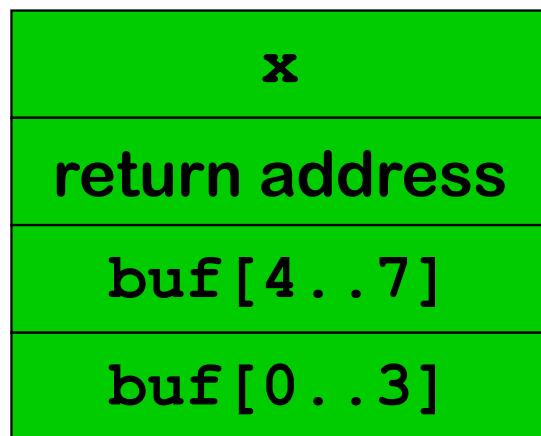
Security canaries are **runtime protection mechanisms** that detect **stack buffer overflows** before they can be exploited.

Security canaries act as **sentinels (canaries) placed in memory** to identify unauthorized modifications, similar to how **canary birds in coal mines** detected toxic gases.

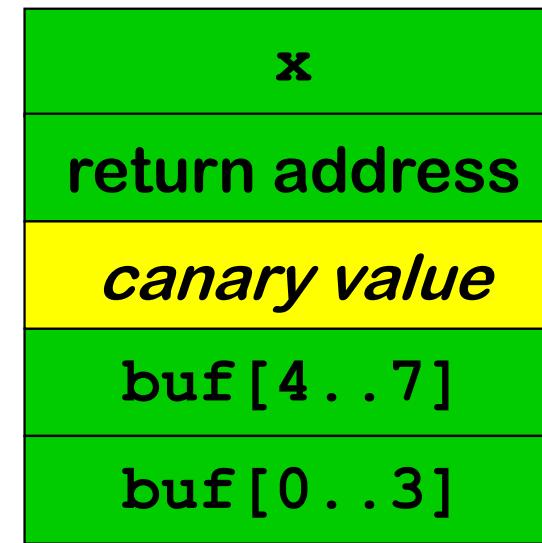
# How canaries detect buffer overflow?

- Canaries are values inserted in the activation record between local variables and the function return address on the stack.
- If a buffer overflow occurs and an attacker overwrites adjacent memory, the canary value is modified.
- Before returning from a function, the program checks if the canary has changed.
  - If the value is unchanged, execution continues.
  - If the value is modified, the program terminates immediately, preventing exploitation.

## Stack without canary



## Stack with canary



# What is a canary value?

**Canaries are Random Values**

- Generated randomly at program startup.

Hard for attackers to predict.

# Canaries in the compilation process

---

Canaries are run-time data structures to prevent **buffer overflow attacks**

---

The integration of **stack canaries** into the compilation process is a security enhancement.

---

The **compiler instruments the code to generate, insert, and verify canary values** at runtime.



## Compilation steps

### Canary Generation (Program Startup)

When the program starts executing, a random canary value is generated.

This value is stored in a global memory location (e.g., fs:0x28 on Linux for x86-64 systems).

Compilation steps

Canary Insertion  
(Prologue)

---

The compiler **modifies function entry code (prologue)** to **insert the canary before the return address**.

---

This ensures that **any overflow from local variables must first overwrite the canary** before reaching the return address.

Compilation steps

Canary Insertion  
(Example)

---

```
mov %fs:0x28, %rax      # Load global canary value
```

*At assembler level:*

---

```
mov %rax, -0x8(%rbp)    # Store it onto the stack
```

*Prologue*



## Compilation steps

### Canary Verification (Epilogue)

---

Before returning from a function, the compiler inserts in the ~~prologue~~<sup>epilogue</sup> code a check to see if the canary is unchanged.

---

If the canary value was modified, the program triggers a security alert and aborts execution



## Compilation steps

### Canary Verification (Example)

---

```
mov -0x8(%rbp), %rdx
# Load stored canary
xor %rdx, %fs:0x28
# Compare with original value
jne stack_smash
# If changed, abort execution
```

---

If jne stack\_smash is triggered, the program calls:

**\_\_stack\_chk\_fail();**

This **function terminates execution** to prevent further exploitation.

## Canaries: Discussion

- **Prevents return address overwrites** (common in buffer overflow exploits).
- **Widely deployed in modern systems** (Linux, Windows, macOS).

# Canaries: Weakness

- **Heap overflows & format string attacks** are **not protected**.
- **Bypassing techniques:**
  - **Memory disclosure** (leaks canary value).
  - **ROP (Return-Oriented Programming)** attacks avoid direct return address modifications:
  - Instead of injecting new code, the attacker **chains small snippets of existing code (gadgets)** to execute malicious operations.

## What is ROP?

ROP is a method where an attacker takes control of a program's execution flow without injecting new code. Instead of writing their own shellcode (which would be blocked by modern protections), they reuse small sequences of instructions—called gadgets—that are already present in the program's memory (usually in libraries like libc or the program's own code). Each gadget typically ends in a ret (return) instruction, allowing the attacker to "chain" them together to perform arbitrary actions.

## What are ROP Attacks?

ROP attacks exploit vulnerabilities like buffer overflows to overwrite the stack and control the return address. Instead of pointing to a single attacker-controlled address, they carefully redirect execution through a sequence of gadgets. This allows them to bypass security mechanisms (like non-executable memory) and execute malicious operations without injecting new code.

You're absolutely right! If the canary is intact, the system won't detect an issue when the function returns. But there's still a small detail: even if the attacker leaks the canary and overwrites the return address correctly, the program would still return only once.

## How Does ROP Avoid Canary Checks?

The trick is that ROP doesn't rely on a single return—it chains multiple returns without triggering a function epilogue that checks the canary. Here's how:

# How Attackers Bypass Stack Canaries

If an attacker can leak the canary value, the attacker can overwrite the buffer without modifying the canary, ensuring that the check passes.

# Leaking Canaries: with format strings

```
void vulnerable_function() {  
    char buffer[64];  
    printf("Enter input: ");  
    gets(buffer); // Buffer Overflow risk!  
    printf(buffer); // String vulnerability leaks stack data  
}  
  
int main() {  
    vulnerable_function();  
    return 0;  
}
```

# Leaking Canaries: with format strings

```
void vulnerable_function() {  
    char buffer[64];  
    printf("Enter input: ");  
    gets(buffer); // Buffer Overflow risk!  
    printf(buffer); // String vulnerability leaks stack data  
}  
  
int main() {  
    vulnerable_function();  
    return 0;  
}
```

The attacker input:  
%p-%p-%p-%p-%p

# Leaking Canaries: with format strings

```
void vulnerable_function() {  
    char buffer[64];  
    printf("Enter input: ");  
    gets(buffer); // Buffer Overflow risk!  
    printf(buffer); // String vulnerability leaks stack data  
}  
  
int main() {  
    vulnerable_function();  
    return 0;  
}
```

The attacker input:  
%p-%p-%p-%p-%p

# Leaking Canaries: with format strings

```
void vulnerable_function() {  
    char buffer[64];  
    printf("Enter input: ");  
    gets(buffer); // Buffer Overflow risk!  
    printf(buffer); // String vulnerability leaks stack data  
}  
  
int main() {  
    vulnerable_function();  
    return 0;  
}
```

The attacker input:  
%p-%p-%p-%p-%p

%p-%p-%p-%p-%p?  
the format specifier %p prints **memory addresses**.  
The program executed uses printf() with %p-%p-%p-%p-%p to print out stack values, including the **stack canary**

%p: print a memory address. I will display the structure of the stack and get access to it

## Another example

```
void secret_function() {  
    printf("You've bypassed the stack canary!\n");  
    system("/bin/sh"); // Spawns a shell  
}  
  
void vulnerable_function() {  
    char buffer[64];  
    printf("Enter input: ");  
    gets(buffer); // Buffer Overflow risk!  
}  
  
int main() {  
    vulnerable_function();  
    return 0;  
}
```

Attacker knows that the canary is always put right after the return address. But what if I modify a function pointer without touching canary to point? Attacker must know layout of AR.

# Bypassing canaries: predicting memory addresses

- The attacker can predict memory addresses, making bypassing stack canaries easier.
- The issues: the stack canary is always at a predictable memory location.
  - If attackers can leak stack memory (e.g., format string attack), they can directly read the canary value.
  - With the known canary value, attackers can overwrite return addresses without triggering stack protection.

# Example

```
void secret_function() {  
    printf("You've bypassed the stack canary!\n");  
    system("/bin/sh"); // Spawns a shell  
}  
  
void vulnerable_function() {  
    char buffer[64];  
    printf("Enter input: ");  
    gets(buffer); // Buffer Overflow risk!  
}  
  
int main() {  
    vulnerable_function();  
    return 0;  
}
```

**The attack strategy:**  
**overwrite a saved function pointer** near  
the buffer.  
Redirect execution to `secret_function()`,  
bypassing the canary check.

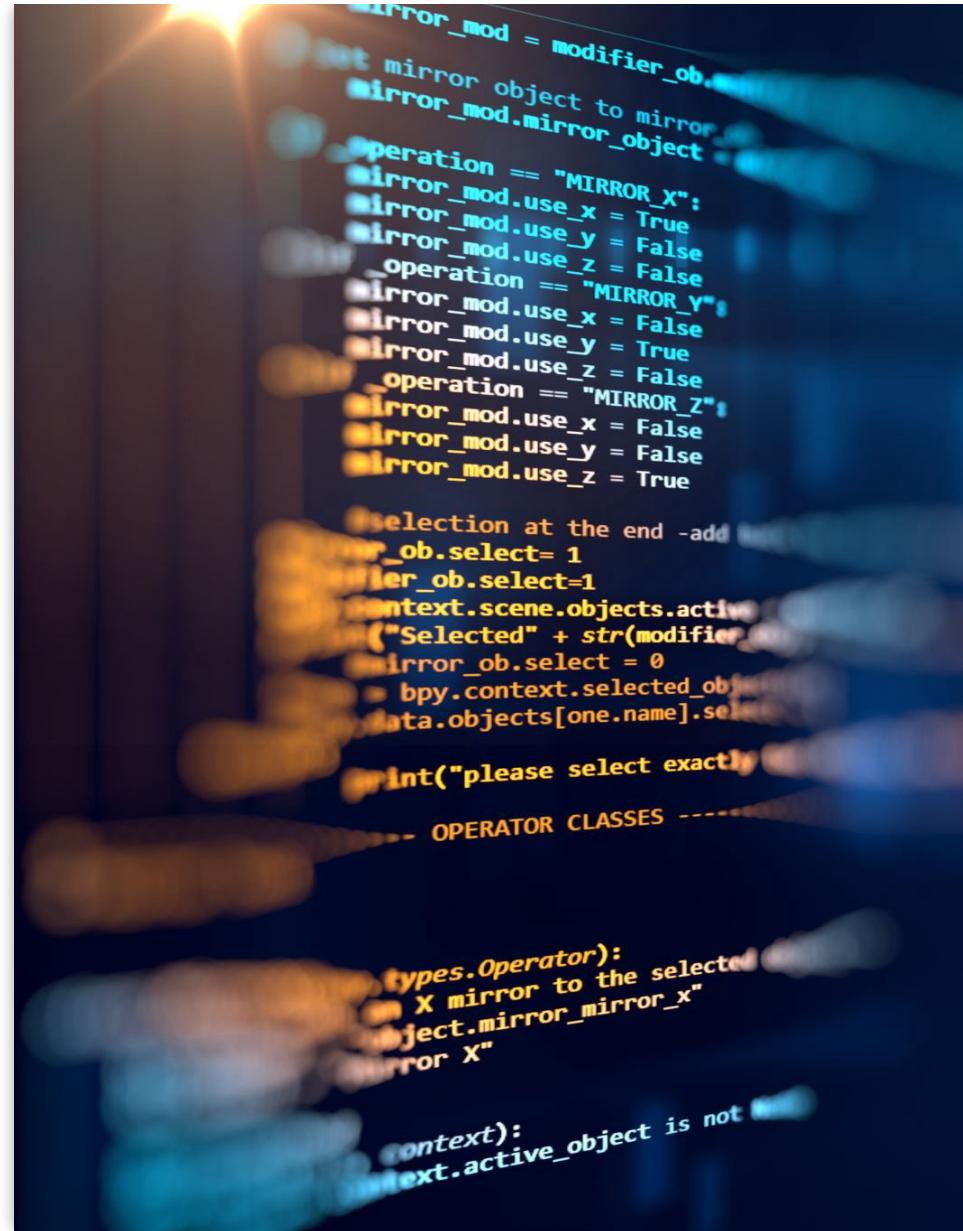
# Address Space Layout Randomization (ASLR)

- **Address Space Layout Randomization (ASLR)** is a security technique used to **randomize memory addresses** of program components **at runtime**.
- This makes it **harder for attackers** to exploit **memory-based vulnerabilities**, such as **buffer overflows** and **return-oriented programming (ROP) attacks**.
- **Key Idea:** If an attacker **doesn't know where key memory structures are located**, they **can't reliably overwrite them** or execute malicious payloads.

# How ASLR Work

Every time a program runs,  
ASLR randomizes the locations of:

- Stack
- Heap
- Shared libraries
- Executable binary
- Memory-mapped files



```
    mirror_mod = modifier_ob.  
    # mirror object to mirror  
    mirror_mod.mirror_object =  
        operation == "MIRROR_X":  
            mirror_mod.use_x = True  
            mirror_mod.use_y = False  
            mirror_mod.use_z = False  
        operation == "MIRROR_Y":  
            mirror_mod.use_x = False  
            mirror_mod.use_y = True  
            mirror_mod.use_z = False  
        operation == "MIRROR_Z":  
            mirror_mod.use_x = False  
            mirror_mod.use_y = False  
            mirror_mod.use_z = True  
  
    # selection at the end - add  
    _ob.select= 1  
    mirr_ob.select=1  
    context.scene.objects.active =  
        ("Selected" + str(modifier))  
    mirror_ob.select = 0  
    bpy.context.selected_objects =  
        data.objects[one.name].select  
  
    print("please select exactly one object")  
  
- OPERATOR CLASSES -  
  
@types.Operator:  
    X mirror to the selected object.mirror_mirror_x"  
    mirror X"  
  
@context:  
    context.active_object is not None
```

**0x400000 - Executable Code**

**0x600000 - Data Segment**

**0x7fff0000 - Stack**

**0x08048000 - Heap**

## Predictable Memory Layout

**0x45d20000 - Executable Code**

**0x6af50000 - Data Segment**

**0x77bf9000 - Stack**

**0x08e34000 - Heap**

## Randomized Memory Layout

*Today we work with them all randomized*

# Address Space Layout Randomization

- Attacker needs detailed info about memory layout
  - to jump to specific piece of code
  - to corrupt a pointer at known position on the stack
- Attacks become harder if compilers randomise the memory layout every time we start a program
  - change the offset of the heap, stack, etc, in memory by some random value
- Attackers can still analyse memory layout on their own machine, but will have to determine the offsets used on the victim's machine to carry out an attack

# ASLR Benefits

- Prevents Return-Oriented Programming (ROP) Attacks
  - ROP attacks rely on **fixed addresses of instructions** to execute arbitrary code. With ASLR, attackers **cannot reliably predict gadget locations**.
- Mitigates Buffer Overflow Exploits
  - Overwriting a return address **without knowing its exact location** causes **crashes instead of successful exploitation**.
- Prevents Memory Disclosure Attacks
  - Even if an attacker leaks memory addresses, **they only work for the current execution**.
  - **Next execution = new random addresses.**

# Shadow Stack

Extension of the run-time data structure (stack) in the abstract machine  
Only used to store return addresses

# Mitigating Control Flow Attacks

Control flow attacks, such as stack-based buffer overflows, exploit vulnerabilities in the program's control flow to execute malicious code.

A **shadow stack** is a security mechanism designed to counter such attacks by maintaining an independent, protected copy of return addresses.

# How the Shadow Stack Works

---

A shadow stack is a run time data structure that operates as follows:

---

**Call Instruction:** When a function is called, the return address is pushed onto both the normal stack and the shadow stack.

---

**Return Instruction:** Before executing a return instruction, the system verifies that the return address on the normal stack matches the corresponding entry in the shadow stack.

---

**Mismatch Detection:** If the return addresses do not match, it indicates a possible control flow attack, and the program is terminated or other security measures are triggered.

---

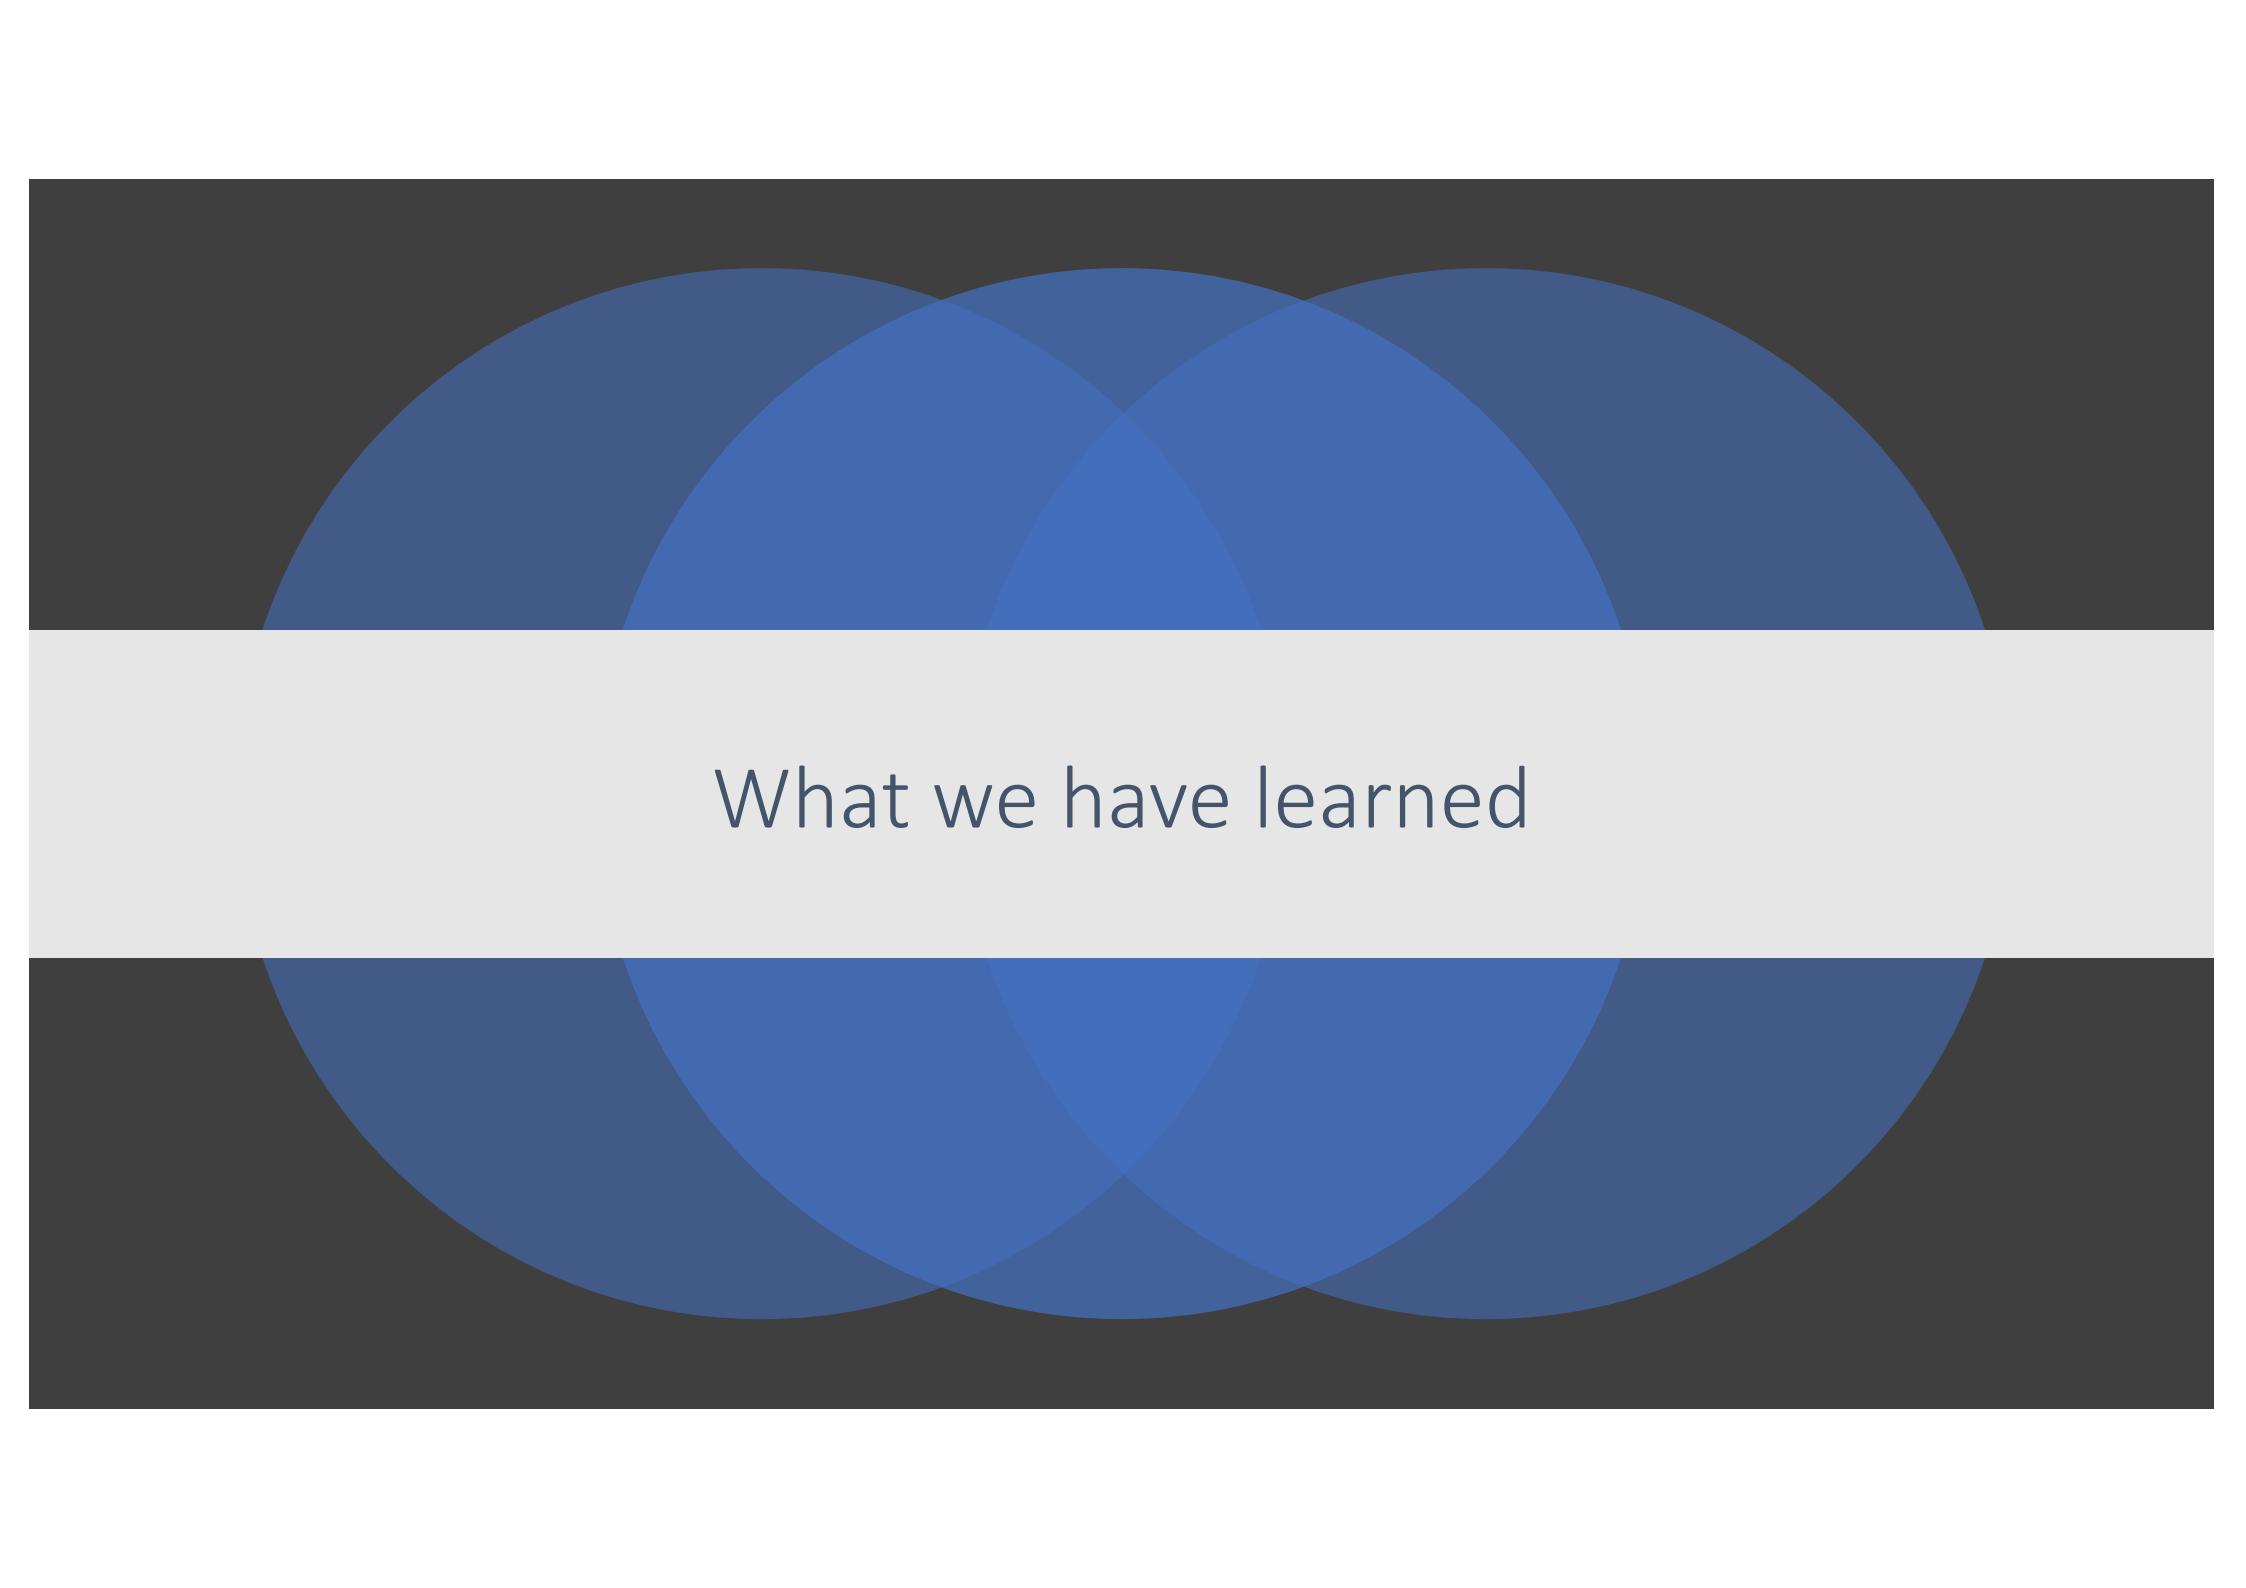
**Protected Memory Region:** The shadow stack is stored in a protected memory region that prevents unauthorized writes, ensuring its integrity.

# Discussion

- **Performance Overhead:** Maintaining a separate stack and performing additional checks introduce some runtime overhead.
- **Compatibility Issues:** Legacy applications and certain optimization techniques may not be fully compatible with shadow stacks.
- **Bypass Techniques:** Sophisticated attackers may attempt to locate and tamper with the shadow stack, requiring further protections like hardware-based enforcement.

# Discussion

- Intel Control-Flow Enforcement Technology (CET): Includes shadow stack support to protect return addresses at the hardware level.  
*↑ Protection at HW level for intel family*
- Compiler-Based Approaches: Certain compilers (e.g., LLVM) implement software-based shadow stacks for enhanced security.



What we have learned

Run-time  
organization  
and code  
generation

## Stack & Heap

- Stack Canaries
- ASLR
- Shadow stack

## Code instrumentation

↳ generated by compiler should be written  
to use these solutions

Impact on compiler and runtime organization

# Non Executable Memory

- **Non-executable (NX) memory** is a security feature that prevents certain memory regions from executing code, thereby mitigating exploits like **buffer overflows** and **code injection attacks**.
- Implemented using hardware-supports.

# How NX works

- **Memory Segmentation:** The operating system marks specific memory regions (e.g., stack, heap) as non-executable.
- **Execution Prevention:** If an attacker injects and tries to execute malicious code in these protected regions, the CPU generates a segmentation fault or access violation.
- **Legitimate Execution Allowed:** Code execution is restricted to designated areas like text (code) segments, while writable regions remain non-executable.

# Discussion

---

The interpreter refuses to execute non-executable code

---

Attackers can then no longer jump to their own attack code, as any input provided as attack code will be non-executable

---

Enhances Overall System Security: Reduces the attack surface by enforcing strict execution policies.

Regions now are strict.

But ...

NX memory does not work for JIT (Just In Time) compilation, where e.g. JavaScript is compiled to machine code at run time.

## **Just-In-Time (JIT) Compilation**

- JIT compilation is a technique where code is compiled at runtime instead of ahead of time. This allows languages like JavaScript or Java to convert high-level code into machine code just before execution rather than being pre-compiled.

### **The Conflict**

The issue arises because JIT compilation requires the ability to execute code in places that might be marked as non-executable. For example:

- When a JavaScript engine (like V8 in Chrome or SpiderMonkey in Firefox) executes JavaScript code, it often converts the script into machine code dynamically using JIT compilation for performance.
- This dynamically generated machine code needs to be executed in memory. If the memory where this machine code is generated is non-executable, the system will prevent it from running, causing an issue for JIT-based environments.

## **2. Performance Considerations**

- JIT compilers are designed to run code on-the-fly, optimizing for performance and responsiveness. This means they typically generate machine code in memory locations that are closer to the execution context (e.g., within the runtime memory of a web page, application, or script).
- Relocating this code to predefined, protected executable sections would introduce overhead and might reduce the performance benefit of JIT, because JIT compilation is meant to be fast and flexible. It would require extra steps to manage where and how the code is generated and executed.



Enclaves

# Security enclaves

HW trusted execution environment

A security enclave is a protected execution environment designed to isolate sensitive computations and data from the rest of the system.

Enclaves isolates part of the code together with its data

- Code outside the enclave cannot access the enclave's data
- Code outside the enclave can only jump to valid entry points for code inside the enclave

Security enclaves ensure confidentiality, integrity, and code authenticity, making them crucial for building trusted execution environments.

Code outside the island cannot assess the enclave.

You enclose code and data in an island. Code of attacker is outside the enclave, so attacker cannot run what's in the enclave

## Example: security-sensitive code in larger program

secret.c

```
static int tries_left = 3;
static int PIN = 1234;
static int secret = 666;

int get_secret (int pin_guess) {
    if (tries_left > 0) {
        if ( PIN == pin_guess ) {
            tries_left = 3; return secret; }
        else {
            tries_left--; return 0 ;}
    } }
```

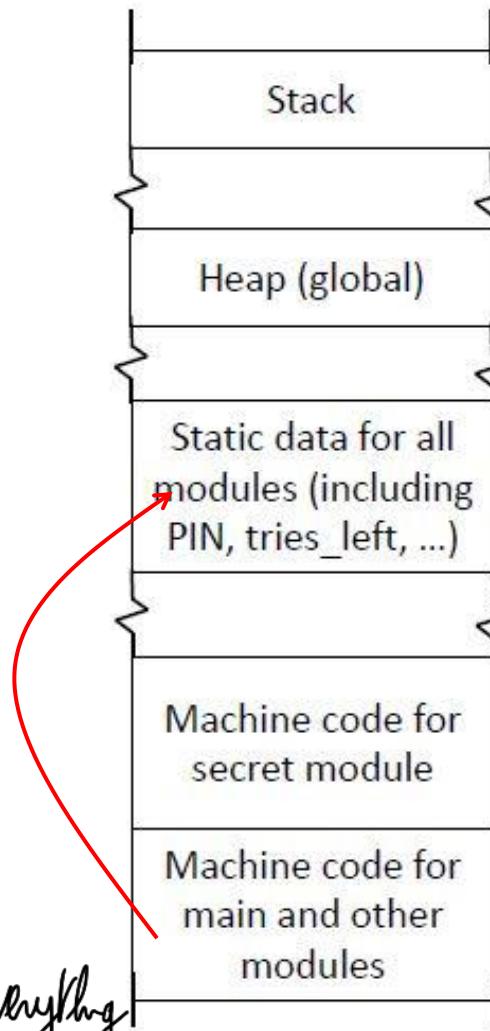
main.c

```
# include "secret.h"
... // other modules
void main () {
...
}
```

Bugs or  
malicious code  
anywhere in the  
program could  
access the  
high-security data

Without isolating everything

is accessible



Example from [N. van Ginkel et al, Towards Safe Enclaves, HotSpot 2016]

## Isolating security-sensitive code with secure enclaves

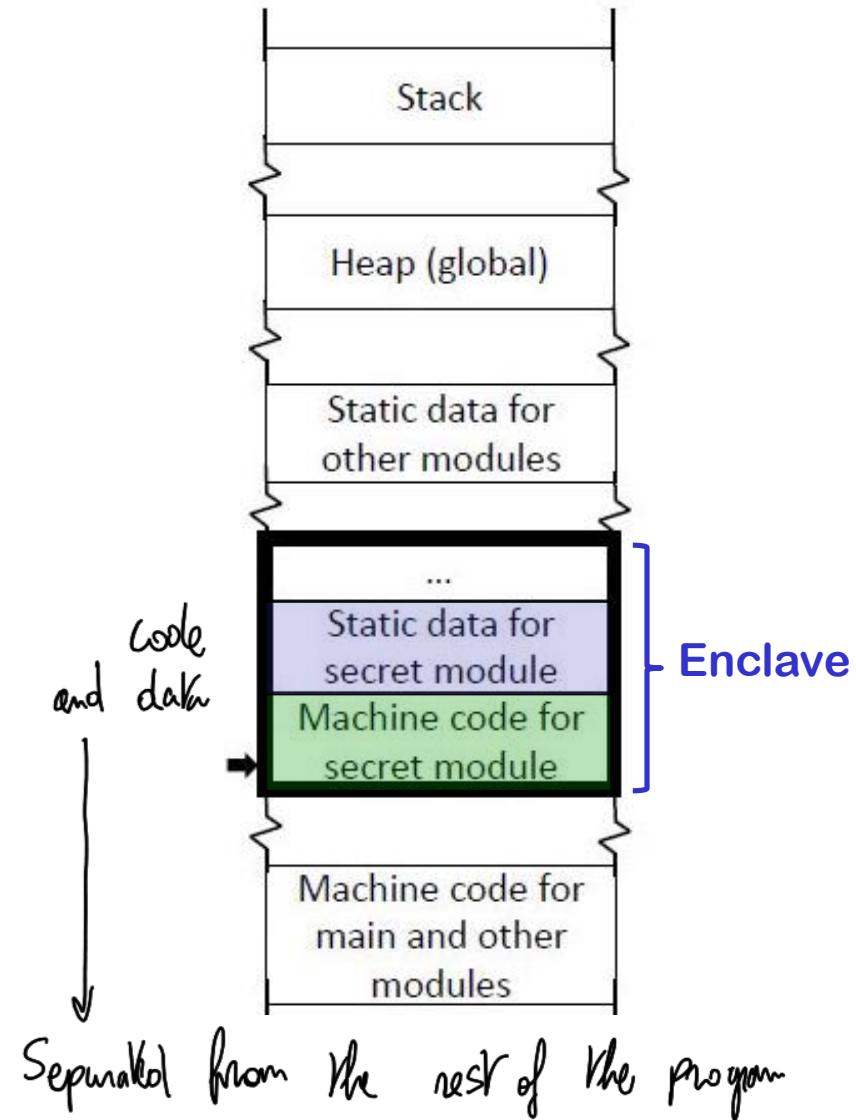
secret.c

```
static int tries_left = 3;
static int PIN = 1234;
static int secret = 666;

int get_secret (int pin_guess) {
    if (tries_left > 0) {
        if ( PIN == pin_guess) {
            tries_left = 3; return secret; }
        else {
            tries_left--; return 0 ;}
    }
}
```

main.c

```
# include "secret.h"
... // other modules
void main () {
...
}
```



## Isolating security-sensitive code with secure enclaves

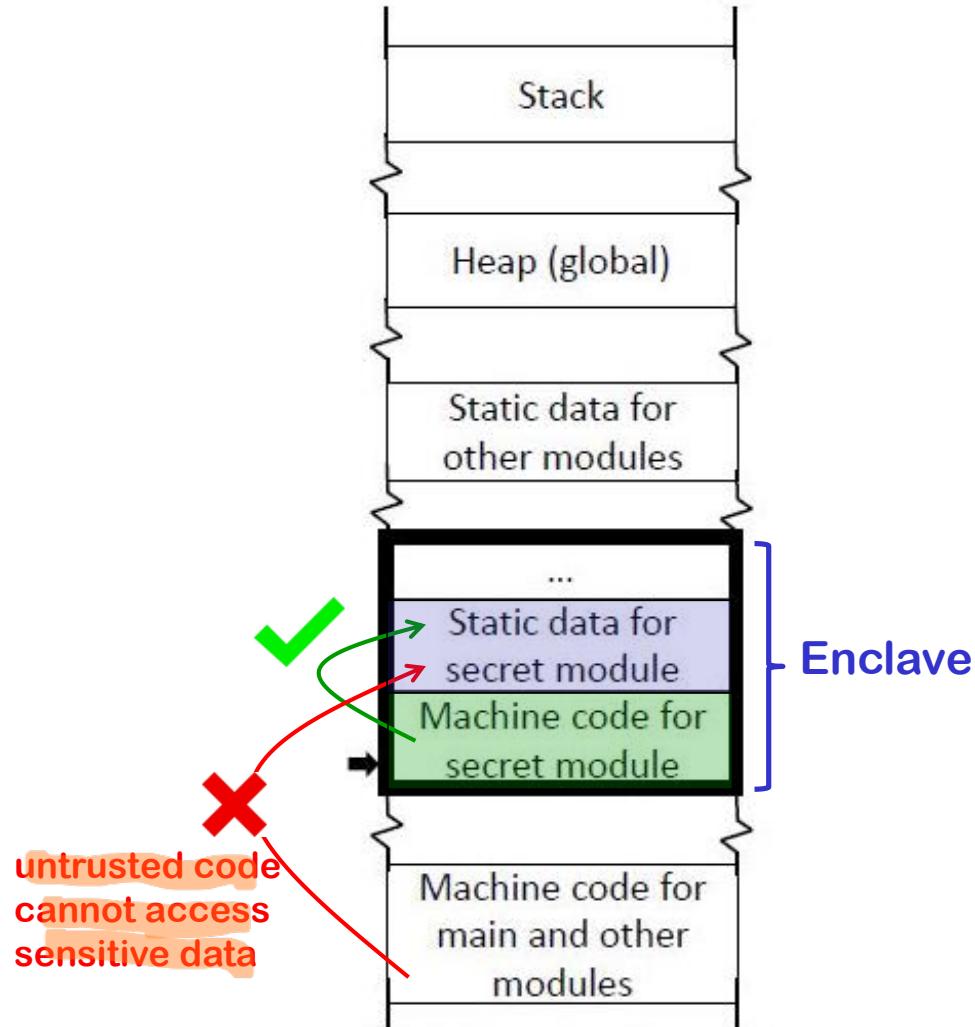
secret.c

```
static int tries_left = 3;
static int PIN = 1234;
static int secret = 666;

int get_secret (int pin_guess) {
    if (tries_left > 0) {
        if (PIN == pin_guess) {
            tries_left = 3; return secret;
        } else {
            tries_left--; return 0;
        }
    }
}
```

main.c

```
# include "secret.h"
... // other modules
void main () {
    ...
}
```



## Isolating security-sensitive code with secure enclaves

secret.c

```
static int tries_left = 3;
static int PIN = 1234;
static int secret = 666;

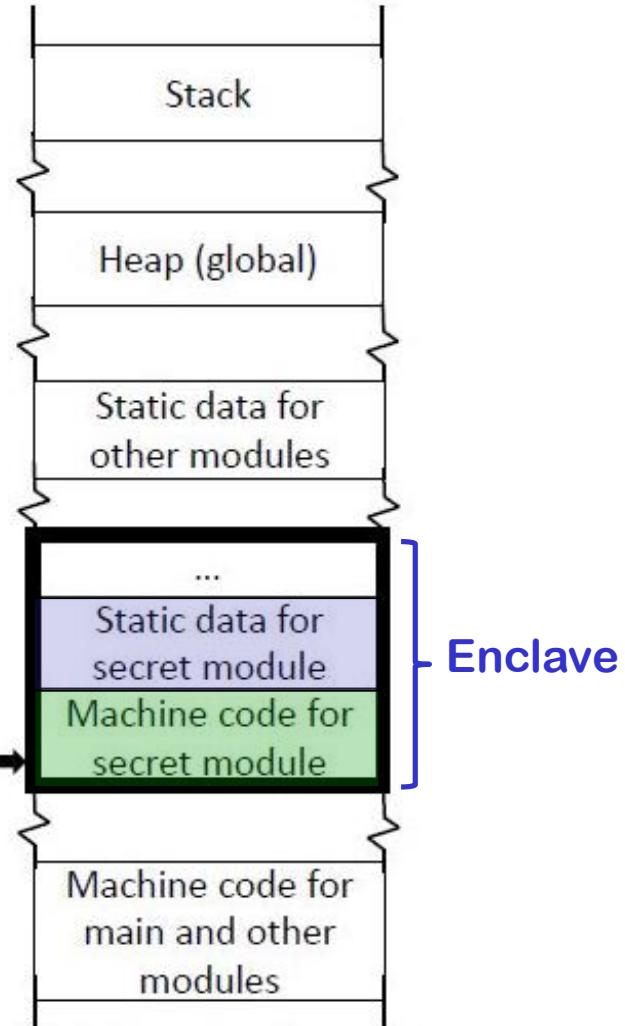
int get_secret (int pin_guess) {
    if (tries_left > 0) {
        if (PIN == pin_guess) {
            tries_left = 3; return secret; }
        else {
            tries_left--; return 0; }
    } }
```

main.c

```
# include "secret.h"
... // other modules
void main () {
    ...
}
```

Only allowed entry point  
(for `get_secret`)

Untrusted code should not be  
able to jump to the middle of  
`get_secret` code (recall return-to-  
libc & ROP attacks)



# How enclave work

- **Isolated Execution:** Code and data inside the enclave are protected from external access, including from privileged software like the OS and hypervisor.
- **Hardware-Based Protection:** Modern CPUs (e.g., Intel SGX) provide dedicated instructions to create and manage enclaves.
- **Safe Memory:** Data in an enclave is stored in encrypted memory, preventing unauthorized access.
- **Attestation:** Enclaves can cryptographically prove their integrity to external parties, ensuring they haven't been tampered with.

↳ a sort of certificate proving their integrity

# Secure enclave: hardware support

- **Intel SGX (Software Guard Extensions)**: Hardware-supported enclaves for isolated execution.
- **AMD SEV (Secure Encrypted Virtualization)**: Protects VMs from host-level threats.
- **ARM TrustZone**: Provides a separate secure execution environment for mobile and IoT devices.
- **Apple Secure Enclave**: A dedicated co-processor for secure biometric authentication and cryptographic operations.

# AMD TrustZone: key features

**Hardware-Enforced Isolation:** The processor is divided into two execution environments:

- **Secure World:** Runs trusted code, handling sensitive operations like cryptographic functions and secure boot.
- **Normal World:** Runs regular applications and the operating system but has restricted access to secure resources.

**Secure Boot & Firmware Protection:** FW prod. to avoid modification

- Ensures that only verified and signed firmware can execute, protecting against unauthorized modifications or malware injections.

**Trusted Applications & Key Management:**

- Enables secure execution of applications handling sensitive data, such as payment processing, biometric authentication, and digital rights management (DRM).
- Secure storage and management of cryptographic keys to prevent tampering or extraction.

**Memory and I/O Security:**

- Secure memory regions prevent unauthorized access to critical data, even if the normal OS is compromised.
- Protects communication between system components (e.g., CPU, GPU, peripherals).

This is a trusted execution environment

# Run-time & memory organization

## code generation

### Stack, Heap and code area

- Stack Canaries
- Reorder layout of AR and Heap elements (randomization)
- Shadow stack
- NX memory
- Enclave

### Code instrumentation

↳ Gen. of the code has to be modified  
(for ex. to exploit enclaves)

# Code Instrumentation as a Security Mechanism

↑  
use it as SW mech. to protect code  
with respect to certain attacks

- **Code instrumentation** is a security technique that involves modifying a program's code—either at compile-time, runtime, or dynamically—to monitor, analyze, or enforce security policies. (dynamically)
- It is widely used for detecting vulnerabilities, preventing exploits, and enforcing security policies.

## Code instrumentation examples

- The management of canaries and shadow stack are examples of security instrumentation

# How to instrument pointer?

QUITE TRICKY!!!

Which is the weakness of pointers? With pointers anything you can access  
everything

# Cyclone

Cyclone is a **safe dialect of C** designed to prevent common vulnerabilities like buffer overflows, dangling pointers, and format string attacks.

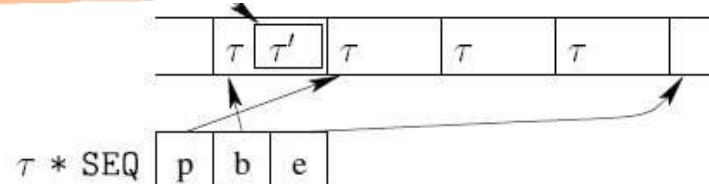
One of its notable features is the use of **fat pointers**, which enhance memory safety while maintaining efficiency.

This is SW instrumentation  
Rust explores the fat pointers idea

level of C prog. language, but  
↑ in a more secure approach

# FAT POINTERS

- Idea: change the representation of pointers to carry bounds information
  - “`int *p`” becomes  
`struct st {int *ptr, int *b, int *e};` Region in which ptr lives  
`struct st p;`
  - the field b points to the beginning of the buffer, and e the end of buffer



# MEMORY OPS

- “ $x = *p$ ” becomes

~~if ( $p.ptr < p.b || p.ptr > p.e$ ) abort();~~

~~$x = * (p.ptr);$~~

ptrs characterise the region of a memory

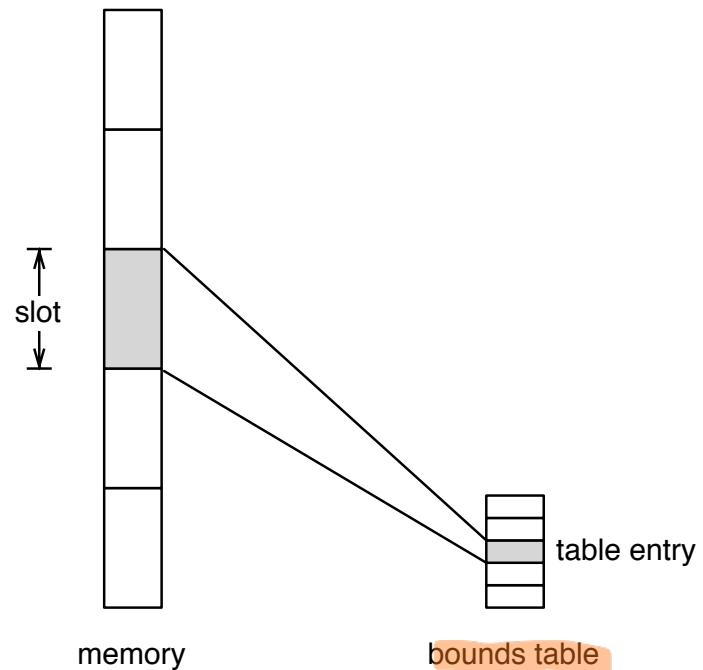
**Bounds Tracking**: any op. has to check the boundaries of a region

- Fat pointers in Cyclone store **both a base address and a length**. This ensures that any access using the pointer stays within **allocated bounds**.
- The compiler and the runtime enforce **boundary checks**, preventing buffer overflows.

# FAT POINTERS: bounds checking

- Compiler-based instrumentation to enforce spatial memory safety for C/C++.
- The general idea is to keep information about all pointers in disjoint metadata, indexed by pointer location

Instrumentation of a ptr  $\rightarrow$  instrument of a comp.  
of compiler. We need a bounds table to  
check if a ptr does something bad.



## **Is This Related to Fat Pointers?**

Yes, but it's **not exactly the same thing**. **Fat pointers** are a specific solution where the **pointer itself** contains the **metadata**, but this slide seems to be **describing an alternative or complementary approach** where the **metadata is separated from the pointer**. Instead of modifying the **pointer itself**, the **compiler tracks metadata about each pointer separately**.

### **2. Indexed by Pointer Location**

- **Metadata** related to a pointer (like its size, bounds, or type information) is **stored separately**, not with the pointer itself.
- This **metadata is "indexed"** by the **location** where the pointer is **stored in memory**.
  - For example, if the pointer **p** is stored at memory location **0x1000**, the **metadata** (like the bounds of the memory region it points to) could be stored in a table where the entry at **index 0x1000 corresponds to p**.
- In simpler terms, when the program uses a pointer (e.g., dereferences it), the system can **look up** the metadata for that pointer by **finding the pointer's location in memory** and retrieving the associated metadata from the metadata table.

# Discussion

- **Preventing Dangling Pointers:** Cyclone's fat pointers integrate with its **region-based memory management**, ensuring that pointers remain valid only within their allocated lifetime.
- **Safe Array and String Operations:** Cyclone replaces **raw C-style arrays and strings** with fat pointers that carry length metadata.
  - Functions operating on strings and arrays use this metadata to prevent unsafe memory access.
- **NULL and Uninitialized Pointer Safety:** Cyclone's fat pointers **cannot be used unless they are properly initialized**.
- **Performance Overhead:** Since fat pointers **include additional metadata**, they consume more memory than traditional C pointers. However, Cyclone optimizes their usage to reduce unnecessary runtime checks when safety can be statically guaranteed.

We gain in memory safety, lose in expressivity (limitation of pointers). Ops for buffers, strings are safe;

Run-time &  
memory  
organization

code  
generation

## Stack, Heap and code area

- Stack Canaries
- Reorder layout of AR and Heap elements (randomization)
- Shadow stack
- NX memory
- Enclaves

## Code Instrumentation

- FAT Pointers

# Security Instrumentation: examples

↳ ability of Java to check at runtime for arrays. In the heap we have memory stored where too,

- **Detecting Buffer Overflows** – Instruments code to track memory accesses and prevent overflows (e.g. Java and JVM)
- **Enforcing Control Flow Integrity (CFI)** – Ensures execution follows a valid control flow path (later). Check if there is an attempt of modifications
- **Runtime Behavior Analysis** – Detects suspicious activity such as unauthorized system calls or memory modifications (later)

(MY) SUMMARY



# Ecosystem

Principles and implementation guidance for software safety and security are well established and broadly accepted.

- Despite these efforts, common types of software defects prevail, and many occupy top ranks of “worst vulnerabilities” lists such as the OWASP Top 10 or the CWE Top 25 Most Dangerous Software Weaknesses for years if not decades.

Common types of defects are introduced during design, development, and deployment they arises from the structure of the *developer ecosystem*

- the end-to-end collection of tools and processes in which developers design, implement, and deploy software. This includes programming languages, software libraries, application frameworks, source repositories, build and deployment tooling and so forth.

SW is part of the ecosystem.  
SW code, PL, tools etc.

# Memory Safety and Security Here

---

Promote the use of memory-safe programming to reduce attack surface.

Promote the adoption of formal methods and related toolkits to detect much earlier rather than tracing back from externally visible systems failures.

A gentle introduction to RUST

A gentle introduction to formal methods