

Dynamic Taint Analysis

Where we are
now?

Security in Programming Languages (Recap)

Run-time Support

- Stack Canaries
- Layout randomization
- Shadow stack
- NX memory
- Enclaves
- FAT Pointer & Data Descriptor
- Web Assembly
- IRM and Code Instrumentation

Static Analysis

- Later

Application Level Reqs: Fine grained control

- Modern apps make security decisions with respect to the abstraction provided by the application level
 - E-Commerce: no goods until payments
 - IoT Apps: avoid data leaks
 - ...



Taking decisions in terms of data:]

Information Security

Valuable (secret) information should not be leaked by computation

Only authorized entities can read from a file (access control approach)

But what happen to information when the entity is authorized?

Idea: what happens to data when auth. is granted?

INFORMATION SECURITY

Typical approach is: ↴

Software systems handles data with different levels of confidentiality and integrity.

Two confidentiality levels,
secret (restricted)
public (unrestricted)

Two integrity levels,
reliable (coming from trusted sources)
dubious (coming from untrusted sources)

INFORMATION FLOW: how the data flow among these levels

Control not only accesses to resources (data at rest)
but also information flow between resources. (data in transit)

Confidentiality policy: information can flow only from less secret to more secret.

"flow"
public → public public → secret
secret → secret secret ↗ public

Integrity policy: information can flow only from more reliable to less reliable.

reliable → reliable reliable → dubious
dubious → dubious dubious ↗ reliable

↑ data is assumed to be reliable

A (STRONGER) MOTIVATION

- Meltdown, Spectre Attacks
- Isolation is a cornerstone of our Trusted Execution Environment (TEE)
 - Modern processors: isolation between the kernel and user processes is realized by a supervisor bit of the processor that defines whether a memory page of the kernel can be accessed or not.
- Meltdown allows overcoming memory isolation by covert channels caused by *out-of-order execution*.
 - Speculative executions

What is speculative execution?

- Speculative execution is an optimization technique used in modern processors. It allows the CPU to guess and execute future instructions before knowing whether they are needed.
 - If the guess is correct, the execution is sped up.
 - If incorrect, the results are discarded, and the correct path is followed.
- This technique is widely implemented in out-of-order execution architectures, where the CPU executes instructions as resources become available rather than strictly following program order.

How speculative execution works

- **Branch Prediction:**
The processor predicts the outcome of a conditional branch (e.g., if-else) before the condition is fully evaluated.
- **Execution of Instructions Speculatively:**
If the processor predicts that a certain branch is likely to execute, it starts computing instructions in that path.
- **Commit or Rollback:**
If the prediction is correct, the results are committed to memory. If incorrect, the processor discards the speculative results and rolls back to the correct path.

Meltdown Attacks

- Meltdown is a hardware vulnerability that exploits speculative execution to leak sensitive data from protected memory regions, such as kernel memory.
- The attack was discovered in 2018 and affects many processors, especially Intel CPUs.

How Meltdown Works

- Meltdown allows an attacker to **read privileged memory** (e.g., kernel space) that **should be inaccessible** to user-level applications.
- It exploits the fact that speculative execution **can access memory that is normally restricted**.

A software view

```
try { i = val;  
      V = PM[i];  
      :  
}  
catch (IllegalReadException)  
{for j in 0 to Max  
  { read PM[j]  
  }  
}
```

A software view

```
try { i = val;      // Attempt to read  
    v = PM[i];    // private data (*)  
    :  
}  
catch (IllegalReadException) {  
    for j in 0 to Max  
        { read PM[j]  
    }  
}
```

Trigger Speculative Execution on Unauthorized Memory

The attacker forces the CPU to execute an instruction that accesses privileged memory (*). Normally, this would cause raising a protection exception, preventing attacker from reading the private data.

A software view

```
try { i = val;      // Attempt to read
      V = PM[i];    // private data (*)
      :              // (speculative execution)
}
catch (IllegalReadException) {
  for j in 0 to Max
    { read PM[j]
  }
}
```

Speculative Execution Reads Restricted Data
Due to **out-of-order execution**, the
CPU speculatively fetches the **restricted data**
before checking access permissions.
This means the data is temporarily loaded into
the **CPU cache**.

A software view

```
try { i = val;      // Attempt to read
      V = PM[i];    // private data (*)
      :
      // (speculative execution)
}
catch (IllegalReadException) {
  for j in 0 to Max
    { read PM[j]
  }
}
```

Side-Channel Attack via Cache

While the CPU eventually realizes the access was illegal and raises the security exception, the cached data remains accessible.

The attacker then uses cache (side-channel attack) to read the leaked data.

Summary

- Meltdown highlights a fundamental weakness in speculative execution: security checks occur too late. While speculative execution boosts performance, it inadvertently exposes memory data that attackers can exploit via side-channel attacks.
- The introduction of suitable hardware fixes helps mitigate the risk, but the trade-off between performance and security remains a challenge for modern processor design.

Meltdown attack

- Meltdown by Moritz Lipp, Michael Schwarz, Daniel Gruss, Thomas Prescher, Werner Haas, Stefan Mangard, Paul Kocher, Daniel Genkin, Yuval Yarom, Mike Hamburg, 2018

Software Mitigation?

- The Meltdown attack exposes vulnerabilities in speculative execution, allowing unauthorized access to privileged memory.
 - Since Meltdown does not rely on traditional software vulnerabilities (e.g., buffer overflows), it is not easily detected by standard security tools.
 - The attack operates at the microarchitectural level, making it invisible to software-level security mechanisms.

Tracking Data Flow to Mitigate Meltdown

↑ approach is track how data flows in execution, and resulted if data under control of the attacker (tainted) is used to make decision, execution is stopped.

Dynamic Taint Analysis (DTA)

Concept: DTA tracks the flow of **tainted data** (e.g., user input) and ensures it does not interact with **sensitive memory**.

Mitigation: If speculative execution tries to **private memory into cache**, **taint analysis** can detect the violation and trigger a response.

What is Dynamic Taint Analysis?

- **Dynamic Taint Analysis (DTA)** is a security technique used in programming languages and runtime environments to track the flow of potentially untrusted data (tainted data) through a program.
- It helps identify vulnerabilities such as **injection attacks, buffer overflows, and unauthorized data leakage**.
- This technique is widely used for detecting and preventing security exploits in languages like **C, C++, Python, Java, and Rust**.

SQL injection is an example in which source of query is under control of attack.

Taint Analysis

- A data-flow analysis mechanism.
 - Purpose: to track the propagations of data.
 - Rule: the variables whose values are computed by exploiting tainted data are also tainted.
- Components:
 - Sources: Taint Seed
 - Untrusted input
 - Sensitive data
 - Propagation: Taint Policy
 - Sinks: Taint Checking/Assert



We have several sources (tainted and untainted) and a lot of mixing. You then have a sink in which data flows to. In the sink you expect only untainted data, otherwise, we stop. So you identify SINK and SOURCES. If you identify a TAINTED PATH, that's bad.

Taint Propagation

- DTA marks specific data (e.g., user input, network data, file input) as tainted.
- When tainted data flows through computations (e.g., string concatenations, arithmetic operations), the taint is propagated to derived values.
- If untainted data combines with tainted data, the result typically remains tainted.



Taint Tracking at Runtime

- Dynamic taint analysis operates at runtime: It tracks taint information across memory, registers, and variables as the program executes. Works at very low level
- This approach ensures precise detection of vulnerabilities but can incur performance overhead.

Taint analysis is an example of an integrity policy.

Policy Enforcement

- The analysis enforces **security policies** to prevent unintended behavior
 - Prevent execution of tainted data in sensitive functions (e.g., SQL queries, shell commands).
 - Block tainted data from reaching unauthorized memory locations.
 - Restrict transmission of tainted data over the network without sanitization.

Fine Grain Taint Tracking

What do we do at level of phys. language? ①. But taint analysis can go even deeper

- Some DTA approaches track taint information at different granularities:
 - **Variable-Level:** Taint is assigned to whole variables. ①
 - **Byte-Level:** Tracks taint at a finer granularity within buffers.
 - **Bit-Level:** Provides the highest precision, though at a higher performance cost.

Taint Analysis for Security

- Identify sources of taint
 - Untrusted input is controlled by the attacker
 - Threat model
- Identify Propagation of taint
 - Tracks flow that manipulates data in order to determine whether the result is tainted or not
- *If the analysis finds that tainted data could be used where untainted data is expected, there could be a potential security vulnerability*



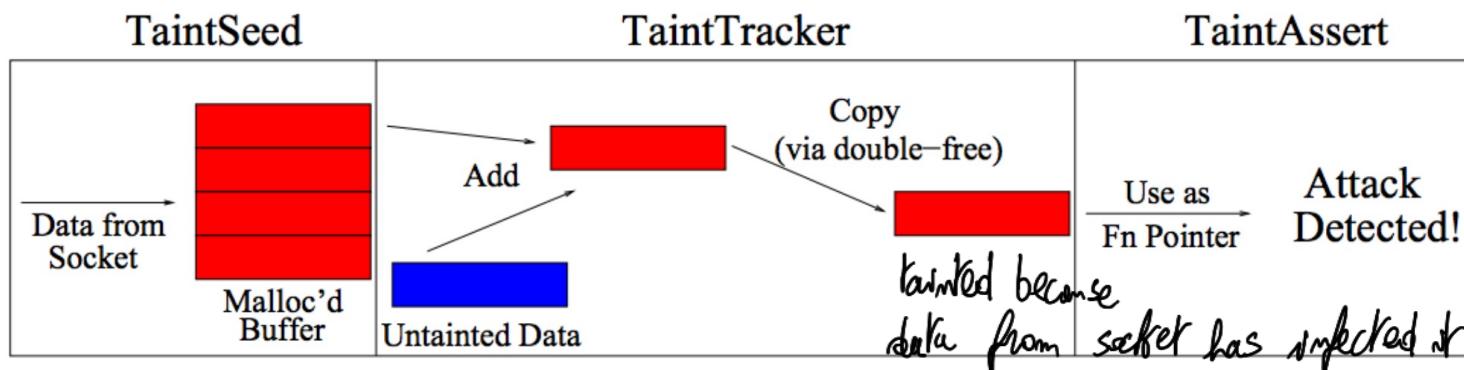
How can DTA be used?

- Tainted data occur in
 - format string attacks.
 - construction of SQL queries, which could be used in SQL injection attacks.
 - Network applications where the attacker controls the communication channels
 - :



Application: Sanity Check

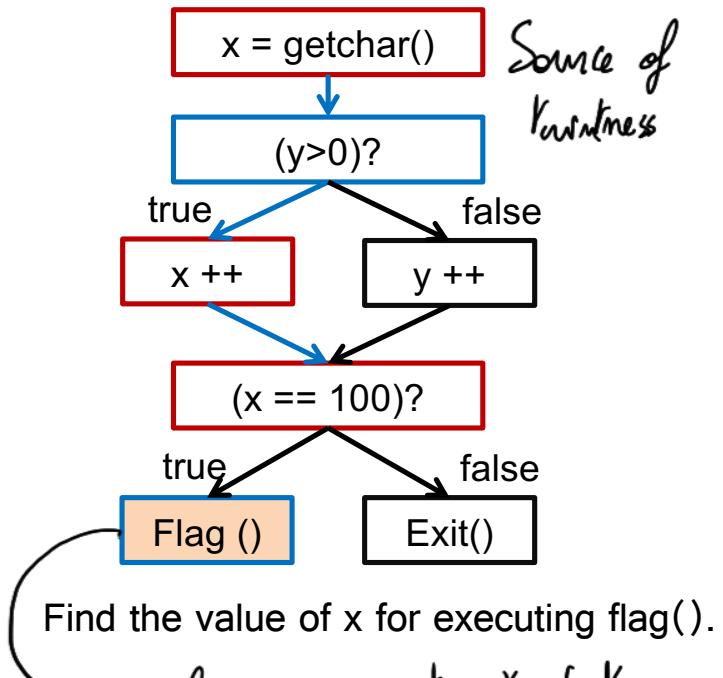
- Check if untrusted data are used dangerously.



James Newsome and Dawn Song, Dynamic Taint Analysis for Automatic Detection, Analysis, and Signature Generation of Exploits on Commodity Software, 2004

Application: Symbolic Execution

Symbolic execution



only requires liveness

- Track the propagation of Variable
- Extract the constraint for executing a particular path.

Sources of data

- Sources of information
- Untrusted – tainted – public
 - Possibly controlled by the attacker
- Trusted – untainted -- secret



MORE INTUITION



tainted



untainted

```

    x = get_input() 
y = x + 42
...
goto y
  
```

Input is tainted

| μ | |
|-------|-----|
| Var | Val |
| x | 7 |

| τ | |
|--------|----------|
| Var | Tainted? |
| x | T |

TaintSeed:

Input $t = \text{IsUntrusted}(\text{src})$
 $\text{get_input}(\text{src}) \downarrow t$

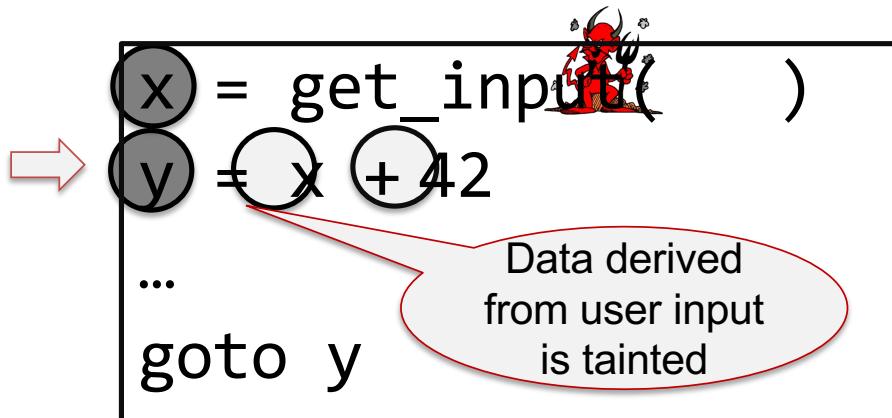
*If source is untrusted
result is untrusted*



tainted



untainted

 μ

| Var | Val |
|-----|-----|
| x | 7 |
| y | 49 |

 τ

| Var | Tainted? |
|-----|----------|
| x | T |
| y | T |

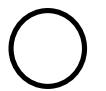
TaintTracker

$$\text{BinOp} \quad \frac{t_1 = \tau[x_1], t_2 = \tau[x_2]}{x_1 + x_2 \downarrow t_1 \vee t_2}$$

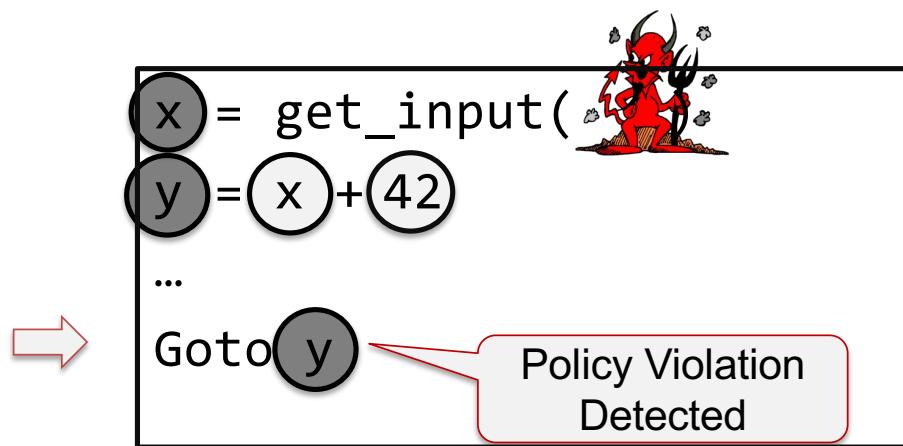
↓ At the level of taintness we take the or of taintness



tainted



untainted



| μ | |
|-------|-----|
| Var | Val |
| x | 7 |
| y | 49 |

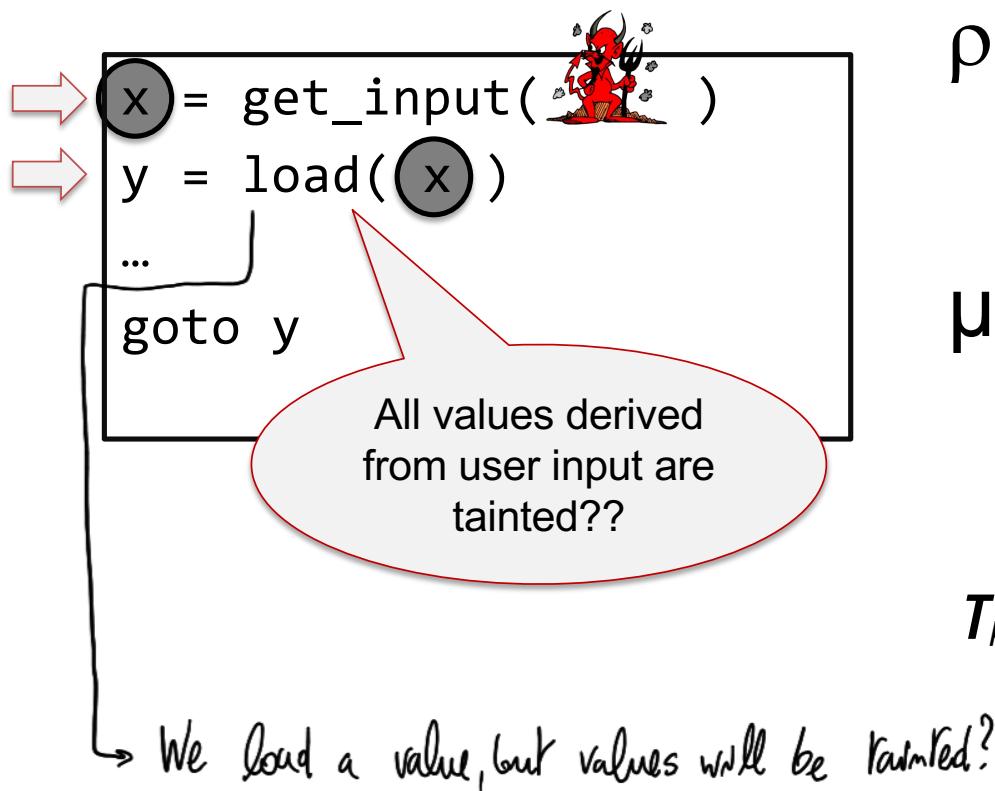
| τ | |
|--------|----------|
| Var | Tainted? |
| x | T |
| y | T |

TaintAssert:

$P_{\text{goto}}(ta) = \neg \text{tainted}$
(Must be true to execute)

Why? Jump/Return address has been overwritten.

Problem: Memory Addresses



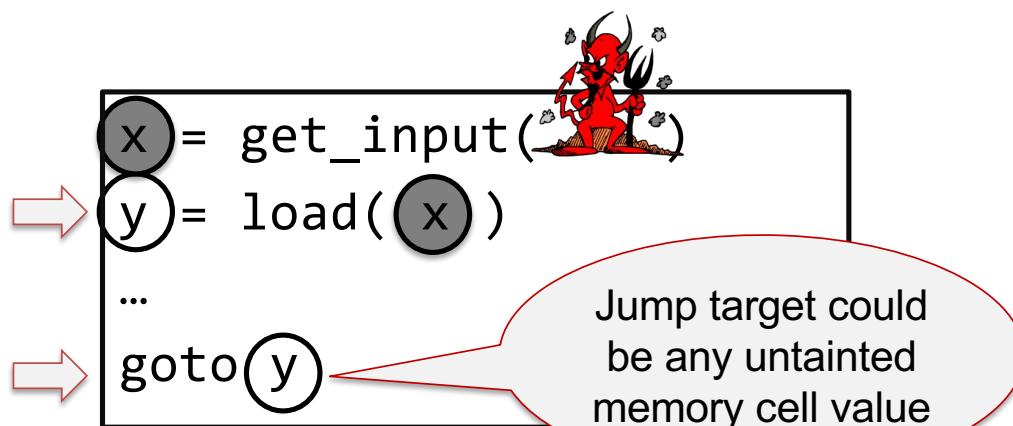
| | Var | Val |
|---------|------|----------|
| ρ | x | 7 |
| | Addr | Val |
| μ | 7 | 42 |
| | Addr | Tainted? |
| T_μ | 7 | F |

Policy 1:

- Propagation of tainted value depends on policy -
It might depend on memory, var., loc.

Taint depends only on the memory cell

$$\text{Load} \quad l = \rho[x], t = \tau_{\mu}[l] \quad * \\ \text{load}(x) \downarrow t$$



| ρ | Var | Val |
|---------|------|----------|
| | x | 7 |
| μ | Addr | Val |
| | 7 | 42 |
| T_μ | Addr | Tainted? |
| | 7 | F |

I associate value associated to x. And find the location. Now it looks in memory and find value.

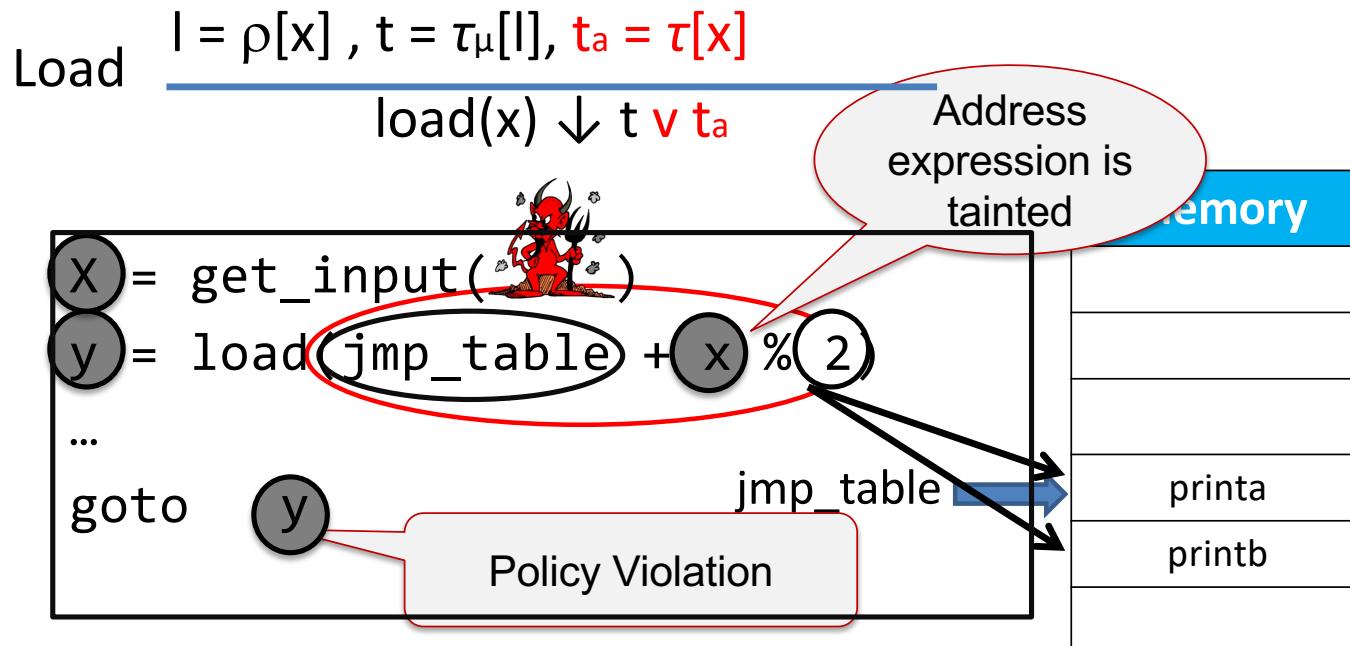
* In premise here you find kind of location. So location is kind.

- ① l vs location in environment associated to x .
- ② What is the rank value associate to location l ?
- ③ γ_M returns association

So rank value of load depends only loc of variable.
We could also take into account environment etc..

Policy 2:

If either the address or the memory cell is tainted, then the value is tainted



The approach

- We explain the main features of dynamic taint analysis by exploiting a simple yet expressive intermediate programming language.
- The language includes features
 - from Java bytecode
 - from assembly languages

An Intermediate programming language (syntax)

```
program    ::=  stmt*
stmt s     ::=  var := exp | store(exp, exp)
               | goto exp | assert exp
               | if exp then goto exp
               | else goto exp
exp e      ::=  load(exp) | exp ◊b exp | ◊u exp
               | var | get_input(src) | v
◊b        ::=  typical binary operators
◊u        ::=  typical unary operators
value v    ::=  32-bit unsigned integer
```

We consider an intermediate prog. language at similar binary level. Bytocode

Remark

The expression
`get_input(src)` returns
input from the source
stream `src`.

↳ only possible
source of input

We model input stream
as a suitable list, ie. `scr =`
`v:: src'`

↑ first el. of input stream and rest of list

We omit the type-
checking
mechanism of our
language and
assume things are
well-typed in the
obvious way,

Operational semantics

Operational semantics of a programming language describes the interpreter: how to execute a program in terms of the run-time values and the run-time data structures

Idea: since dynamic taint analysis is defined in terms of program execution the operational semantics is the natural mechanism on which to base the dynamic taint analysis

First we present op sem., then policy.

Run-time structures

- Σ : the ordered sequence of program statements $\Sigma = \text{Nat} \rightarrow \text{Stmt}$
- μ : memory $\mu: \text{Loc} \rightarrow \text{Values}$ ADDRESSES TO →
IDE entries to locations or to values.
- ρ : environment $\rho: \text{Var} \rightarrow \text{Loc + Values}$ IMMUTABLE VS MUTABLE?
- pc: program counter
- i: next instruction: language has go-to
[nota]

IMMUTABLE: if var goes to loc then value: mutable, otherwise immutable.

MUTABLE:

Program evolution: expressions

The result is our target and witness

$$\mu, \rho \vdash e \Downarrow v$$

Intuition: evaluating the expression e in the run-time context provided by the memory μ and the environment ρ produces v as result

Program evolution: statements

for statements we are interested to
see how machine progresses step
by step

$$\Sigma, \mu, \rho, pc : smt \rightarrow \Sigma, \mu', \rho', pc' : smt'$$

↳ Program is left unchanged (no JIT compilation)

- Intuition: the execution of the statement smt in the run-time context given by
 - the program list (Σ),
 - the current memory state (μ),
 - the current binding for variable (ρ)
 - the current program counter (pc)
- yields a new state of program execution (Σ, μ', ρ', pc')

Remark:
Program evolution-
statements

$$\Sigma, \mu, \rho, pc : smt \rightarrow \Sigma, \mu', \rho', pc' : smt'$$

- Intuition: the execution of the statement smt in the run-time context yields a new state of program execution (Σ, μ', ρ', pc')
- The program Σ does is not modified by transitions.
 - We do not allow programs with dynamically generated code.

A sample of the operational semantics (expressions)



$$\frac{src = v :: src'}{\mu, \rho \vdash \text{getInput}(src) \Downarrow v}$$

expression that produces value v from a list.

$$\frac{\mu, \rho \vdash e \Downarrow v_1 \quad v = \mu(v_1)}{\mu, \rho \vdash \text{load } e \Downarrow v}$$

↳ V dash: we are evaluating this based on M and P.
We resort to memory and environment

$$\frac{}{\mu, \rho \vdash \text{var} \Downarrow \rho(\text{var})}$$

↳ Rule describes behavior.

A sample of the operational semantics (statement)

$$\frac{\mu, \rho \vdash e \Downarrow v \quad \rho' = \rho[\text{var} = v] \quad \iota = \Sigma[\text{pc} + 1]}{\Sigma, \mu, \rho, \text{pc}: \text{var} = e \rightarrow \Sigma, \mu, \rho', \text{pc} + 1: \iota}$$

apply funct. program
↑
old program contin

↓ ↑
We execute this evaluate expression. We have long stop
in the state* Semantics.

* Ev. of exp. is only defined in terms of M, P. While for assignments, it's defined in terms of Σ, M, P, PC .

A sample of the operational semantics (statement)

$$\frac{\mu, \rho \vdash e \Downarrow v \quad \rho' = \rho[\text{var} = v] \quad \iota = \Sigma[pc + 1]}{\Sigma, \mu, \rho, pc: \text{var} = e \rightarrow \Sigma, \mu, \rho', pc + 1: \iota}$$

The current state of
execution

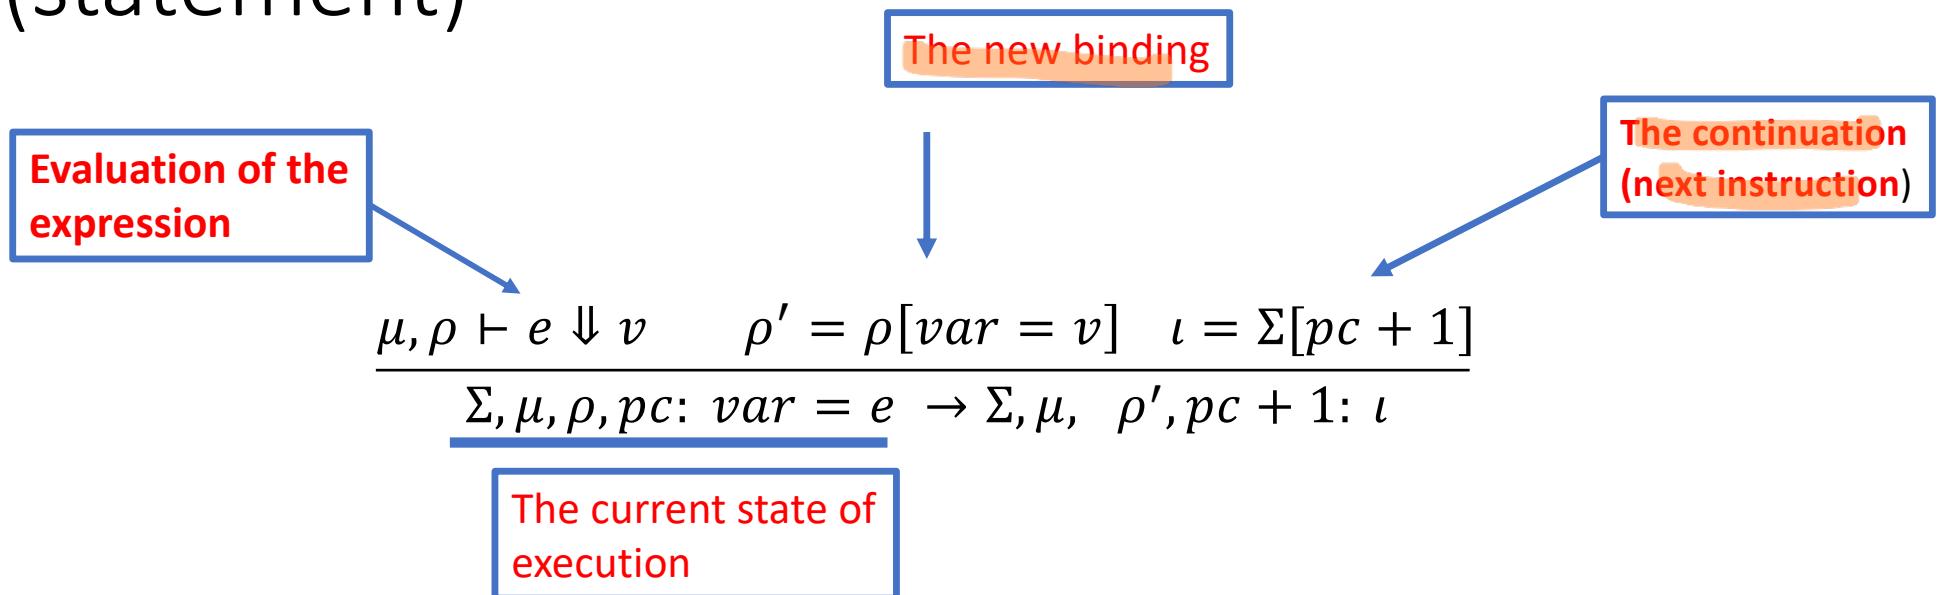
A sample of the operational semantics (statement)

Evaluation of the expression

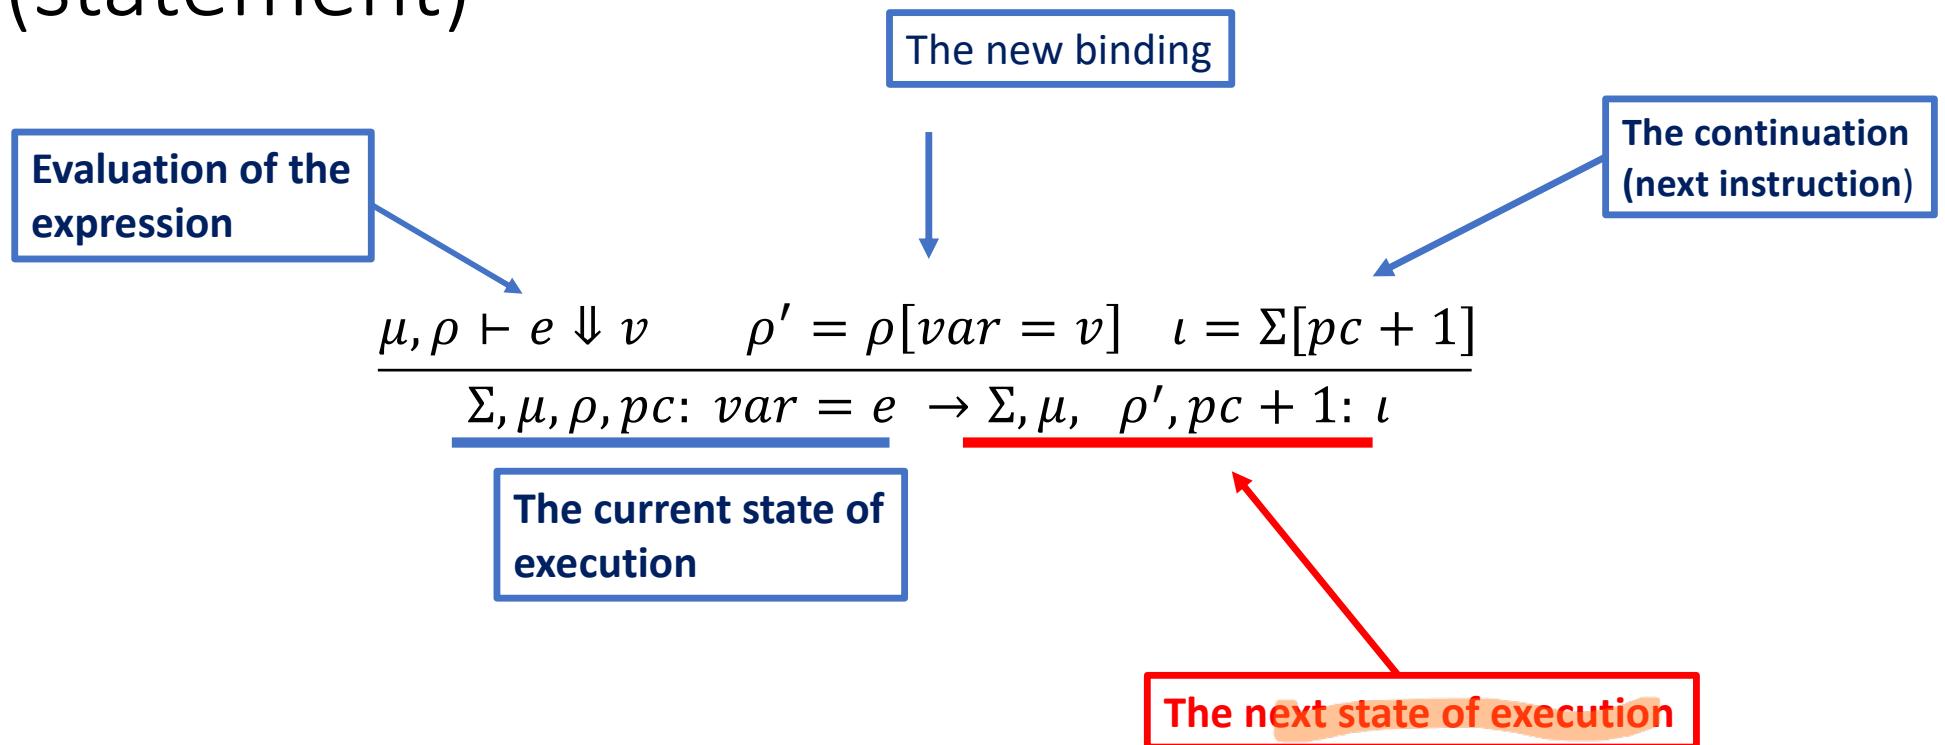
$$\frac{\mu, \rho \vdash e \Downarrow v \quad \rho' = \rho[\text{var} = v] \quad \iota = \Sigma[pc + 1]}{\Sigma, \mu, \rho, pc: \text{var} = e \rightarrow \Sigma, \mu, \rho', pc + 1: \iota}$$

The current state of execution

A sample of the operational semantics (statement)



A sample of the operational semantics (statement)



A sample of the operational semantics (statements)

$$\frac{\mu, \rho \vdash e \Downarrow v_1 \quad \iota = \Sigma[v_1]}{\Sigma, \mu, \rho, pc: goto e \rightarrow \Sigma, \mu, \rho, v_1: \iota}$$

Note structured control flow. You can jump at any possible location

$$\frac{\mu, \rho \vdash e_1 \Downarrow v_1 \quad \mu, \rho \vdash e_2 \Downarrow v_2 \quad \iota = \Sigma[pc + 1] \quad \mu' = \mu[v_1 = v_2]}{\Sigma, \mu, \rho, pc: Store(e_1, e_2) \rightarrow \Sigma, \mu', \rho, pc + 1: \iota}$$

$$\frac{\mu, \rho \vdash e \Downarrow 1 \quad \iota = \Sigma[pc + 1]}{\Sigma, \mu, \rho, pc: assert(e) \rightarrow \Sigma, \mu, \rho, pc + 1: \iota}$$

↑ we check where at a certain point of a program we have something we expect. We find whether certain values have been produced.

```
mirror_mod = modifier_ob
# mirror object to mirror
mirror_mod.mirror_object = ob
operation == "MIRROR_X":
    mirror_mod.use_x = True
    mirror_mod.use_y = False
    mirror_mod.use_z = False
operation == "MIRROR_Y":
    mirror_mod.use_x = False
    mirror_mod.use_y = True
    mirror_mod.use_z = False
operation == "MIRROR_Z":
    mirror_mod.use_x = False
    mirror_mod.use_y = False
    mirror_mod.use_z = True
selection at the end -add
ob.select= 1
ob.select=1
context.scene.objects.active = 
"Selected" + str(modifier)
mirror_ob.select = 0
bpy.context.selected_objects = 
data.objects[one.name].select
int("please select exactly
one object")
- OPERATOR CLASSES ---
```

```
types.Operator):
    X mirror to the selected
    object.mirror_mirror_x"
    or X"
```

```
context):
    context.active_object is not
```

Syntax of our bytecode does not consider
any function code, but we can introduce them

What about here we did in WASM.
Implement prologue, epilogue
functions? and execution. Of course here

**Function calls in
high-level
programming
language are
compiled by storing
the return address
and transferring
control flow.**

We are in a register
based machine

Dynamic taint analysis

Express the taint propagation in terms of the operational semantics of the intermediate language

Dynamic taint analysis is obtained by monitoring the program execution via suitable taint checkers

Design issues

keep track of the taint status of run-time data

Run-time data

$\langle v, \tau \rangle$

Value

Taint status

To manage this representation we need a suitable meta-data and choices to make for design.

We keep data with their own taintness status. This must be reflected at runtime. At runtime data is not just values. We have data + tag.

Run-time structure: extension

Taint $\tau ::= T \mid F$

Trick; Value is just a boolean value

Value ::= $\langle v, \tau \rangle$

τ_p : Maps variable to taint status

τ_μ : Maps addresses to taint status

Environment can map van k immutable data or mem. locations. But k represent k values, we need two more mappings, from variables/addresses to taint values

Taint Policies

We have to identify the source, the sinks (where program expects untainted value) and how tainted value flows in the program

How new taint is introduced

How taint propagates when execution progresses

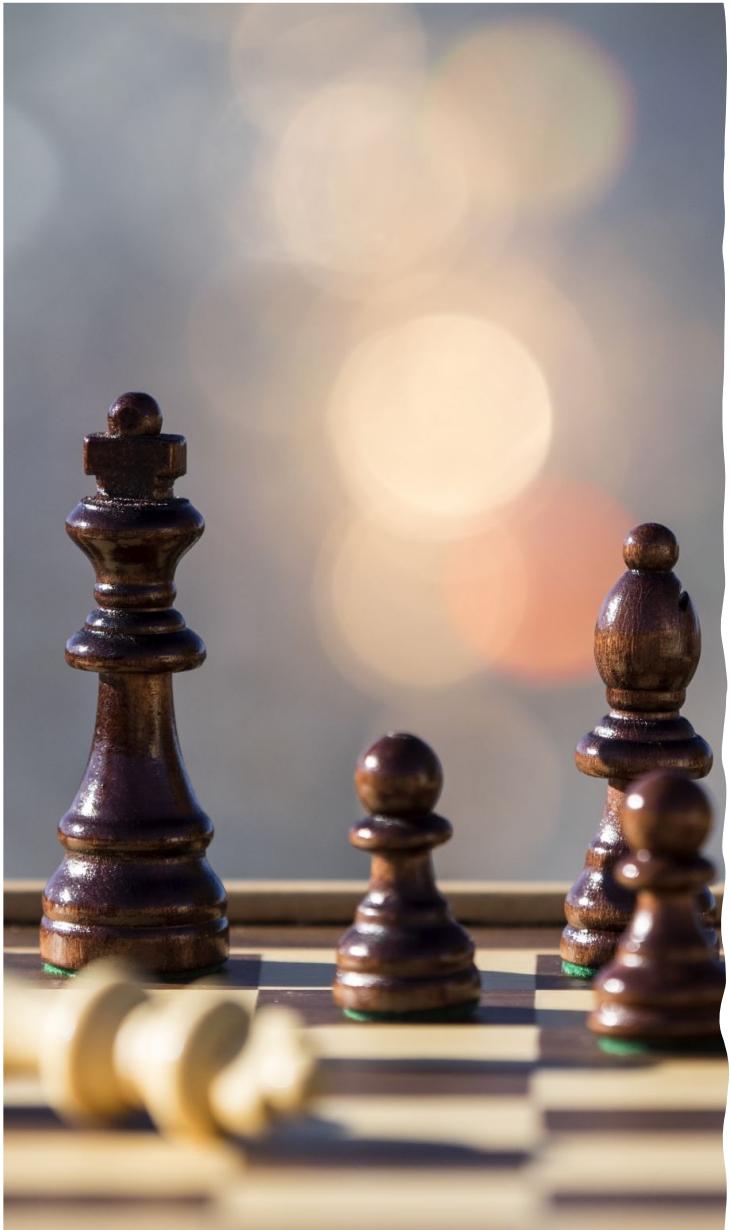
How taint is checked during execution



Taint introducion

Operational rules describe how taint values are introduced in the system

In our language we have a single source: the `getInput` operation.



Taint propagation

The taint propagation rules specifies how taint is derived from operation or control mechanisms

Taint Checking

The taint value of run-time data impacts over the behaviour of programs: the detector may stop the execution if the address of a jump is tainted



Program instrumentation: we perform checking of taint policies before applying execution rules

Taint checking (example)

$$\frac{src = v :: src'}{\tau_\mu, \tau_\rho, \mu, \rho \vdash \text{getInput}(src) \Downarrow \langle v, \text{TPIN}(src) \rangle}$$

Premise: need to find out whether I have at least one value

Extension to keep track of the runtime context for our configuration

TPIN is the taint policy associated to the data source src

TPIN: Taint policy IN

Taint checking (example)

TPIN is the taint policy associated to the data source src

Assume that src is under the control of the attacker

$$\frac{src = 5 :: src'}{\tau_\mu, \tau_\rho, \mu, \rho \vdash getInput(src) \Downarrow \langle 5, T \rangle}$$

A summary of Tainted Policies

| POLICY COMPONENT | POLICY CHECK |
|------------------------------------------------------------------------|----------------|
| TPIN(-) | T |
| TPCONST(-) | F |
| TPASN(t) Assignment: val = exp ⁴ | t |
| TPMEM(t _a t _v) Taint policy takes two values: ① | t _v |
| TPGOTO(t) | not t |
| ① t _a for address, t _v associated | |

To value we want to store in the location. Result is we take the tag associated to value,

A taint status is converted to a boolean value in the natural way, e.g., T maps to true, and F maps to false.

⁴ exp produces a pure value v and tainted t. We take tag associated to value and that will be associated to variable.

GOTO has expression that produces value and result. You cannot jump off value is undefined. You can jump if result value is false. So policy passes when NOT E is verified.

The taint instrumented semantics

$$\frac{src = v :: src'}{\tau_\mu, \tau_\rho, \mu, \rho \vdash getInput(src) \Downarrow \langle v, TPIN(-) \rangle}$$

We apply *Taint* policies

$$\frac{}{\tau_\mu, \tau_\rho, \mu, \rho \vdash v \Downarrow \langle v, TPCONST(-) \rangle}$$

Constant is *untainted*

$$\frac{}{\tau_\mu, \tau_\rho, \mu, \rho \vdash var \Downarrow \langle \rho(var), \tau_\rho(var) \rangle}$$

Lookup of value and *taint value*

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash e \Downarrow \langle v, t \rangle \quad \rho' = \rho[\text{var} = v] \quad \tau'_\rho = \tau_\rho[\text{var} = \text{TPASN}(t)] \quad \iota = \Sigma[pc + 1]}{\tau_\mu, \tau_\rho, \Sigma, \mu, \rho, pc; \text{ var} = e \rightarrow \tau_\mu, \tau'_\rho, \Sigma, \mu, \rho', pc + 1 : \iota}$$

↳ ↑
current states of env and mem

$$\text{TPASN}(t) = t$$

don't you think about your curly hair
clothes that don't go well for you anyway

$$\tau_\mu, \tau_\rho, \mu, \rho \vdash e_1 \Downarrow \langle v_1, t_1 \rangle. \quad \tau_\mu, \tau_\rho, \mu, \rho \vdash e_2 \Downarrow \langle v_2, t_2 \rangle$$

$$\frac{\mu' = \mu[v_1 = v_2] \quad \tau'_\mu = \tau_\mu[v_1 = TPMEM(t_1, t_2)] \quad \iota = \Sigma[pc + 1]}{\tau_\mu, \tau_\rho, \Sigma, \mu, \rho, pc; store(e_1, e_2) \rightarrow \tau'_\mu, \tau_\rho, \Sigma, \mu', \rho, pc + 1: \iota}$$

↑ location ↑ value to store ↑ we won't consider valid states for this location

$$TPMEM(t_1, t_2) = t_2$$



$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash e \Downarrow \langle v, t \rangle \quad TPgoto(t) = T \quad \iota = \Sigma[v]}{\tau_\mu, \tau_\rho, \Sigma, \mu, \rho, pc; goto e \rightarrow \tau_\mu, \tau_\rho, \Sigma, \mu, \rho, v: \iota}$$

when Taint policy produces a true value (so t is false)

TPGOTO(t) = NOT t

The rule is applied when it is safe to perform a jump operation

This holds when $TPgoto(t)$ returns T , i.e. the value t is F (untainted)

When the target address is tainted, $TPgoto(t)$ returns F and the premises of the rule is not satisfied and an exception is raised

1. $x = 2 * \text{get_input}([20]);$
2. $y = 5 + x;$
3. goto y

You'll grow up and
grow so rough and charm them

$$\frac{\text{src} = 20 :: []}{\tau_\mu, \tau_\rho, \mu, \rho \vdash \text{getInput}([20]) \Downarrow \langle 20, T \rangle} \xrightarrow{\text{unboxed}}$$

| | |
|-------------|----|
| τ_ρ | {} |
| ρ | {} |

$x = 2 * \text{get_input}([20]);$

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash 2 * \text{getInput}([20]) \Downarrow \langle 40, T \rangle \quad \rho' = [x = 40] \quad \tau'_\rho = [x = T] \quad \iota = 2}{\tau_\mu, \tau_\rho, \Sigma, \mu, \rho, pc; x = 2 * \text{getInput}([20]) \rightarrow \tau_\mu, \tau'_\rho, \Sigma, \mu, \rho', 2: \iota}$$

| | |
|-------------|----------|
| τ_ρ | $[x=T]$ |
| ρ | $[x=40]$ |

$y = 5 + x;$

$$\frac{\begin{array}{l} \text{Initial status of } x \text{ vs } T \text{ and we summed it with unboxed val.} \\ \tau_\mu, \tau_\rho, \mu, \rho \vdash 5 + x \Downarrow \langle 45, T \rangle \quad \rho' = [x = 40, = 45] \quad \tau'_\rho = [x = T, y = \cancel{x}T] \quad \iota = 3 \end{array}}{\tau_\mu, \tau_\rho, \Sigma, \mu, \rho, pc; y = 5 + x \rightarrow \tau_\mu, \tau'_\rho, \Sigma, \mu, \rho', 3: \iota} \xrightarrow{\text{Result vs Unboxed}}$$

| | |
|-------------|---------------------|
| τ_ρ | $[x=T,$ $y=T]$ |
| ρ | $[x=40,$ $y=45]$ |

goto y

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash y \Downarrow \langle 45, T \rangle \quad F = T \quad \iota = \Sigma[45]}{\tau_\mu, \tau_\rho, \Sigma, \mu, \rho, pc; \text{goto } y \rightarrow err} \downarrow \text{Error}$$

1. $x = 2 * \text{get_input}([20]);$
2. $y = 5 + x;$
3. goto y

Value of PC

| Line # | Stm | ρ | τ_ρ | Rule | pc |
|--------|------------------------------------|----------------|--------------|--------|-----|
| | start | {} | {} | | 1 |
| 1 | $x = 2 * \text{getInput}(20::[]))$ | {x=40} | {x = T} | ASSIGN | 2 |
| 2 | $y = 5 + x$ | {x=40, y = 45} | {x=T, y = T} | ASSIGN | 3 |
| 3 | goto y | {x=40, y = 45} | {x=T, y = T} | GOTO | err |

We didn't analyze loading operation! We have several design choices to manage loading

Load: take #1

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash e \Downarrow \langle v, t \rangle}{\tau_\mu, \tau_\rho, \mu, \rho \vdash \text{load } e \Downarrow \langle \mu(v), \text{TPmem}(t, \tau_\mu(v)) \rangle}$$

↓ Taint value here

↑ location

$$\text{TPmem}(t_a, t_v) = t_v$$

Only the tainted value of the cell is considered

Taint status is: we have taint status of the address and the taint status of the value in that location.

Here we only take into account the taint status associated to that value in the memory.

Problem: we are ignoring the fact that address can be tainted. Who knows what might happen. We accept reading with tainted data.

A main design issue

- A. `x = get_input()`
- B. `y = load(z+x)`
- C. `goto y`

Assumptions:

- The value associated to variable z has been already defined
- The attacker provides input x to the program that is used as a table index offset
- The attacker can provide an appropriate value of x to address any value in memory that is untainted.
- The result of the table lookup is used as the target address for a jump.

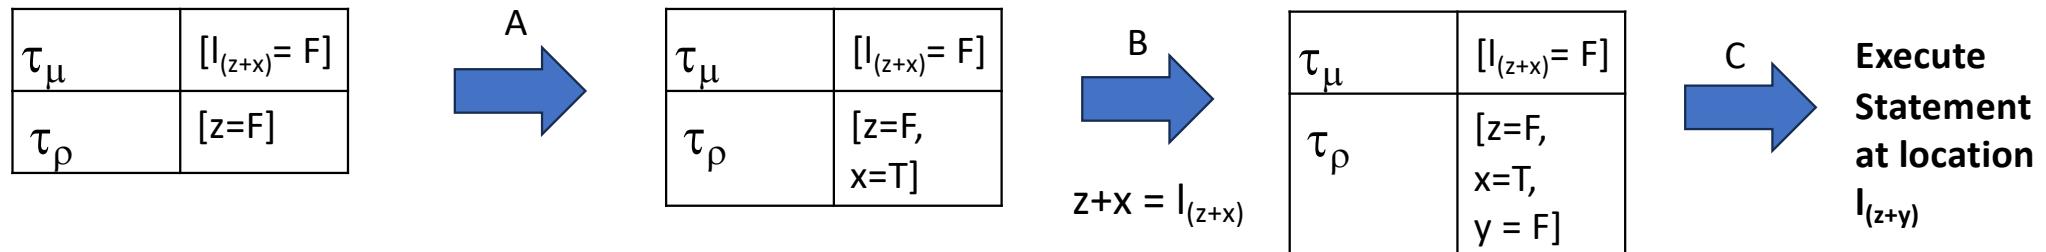
Typical example of RETURN ORIENTED PROGRAMMING

Tainted jump

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash e \Downarrow \langle v, t \rangle}{\tau_\mu, \tau_\rho, \mu, \rho \vdash \text{load } e \Downarrow \langle \mu(v), \text{TPmem}(t, \tau_\mu(v)) \rangle}$$

A. `x = get_input()`
 B. `y = load(z+x)`
 C. `goto y`

$$\text{TPmem}(t_a, t_v) = t_v$$



$$\text{load}(x + y) \Downarrow \langle l_{\{z+x\}}, T \rangle$$

$$\text{TPmem}(T, \tau_\mu(I_{(z+x)})) = \tau_\mu(I_{(z+x)}) = F$$

X

Load: take #2

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash e \Downarrow \langle v, t \rangle}{\tau_\mu, \tau_\rho, \mu, \rho \vdash \text{load } e \Downarrow \langle \mu(v), \text{TPmem}(t, \tau_\mu(v)) \rangle}$$

$\text{TPmem}(t_a, t_v) = t_a \text{ OR } t_v$ Stronger policy for the memory

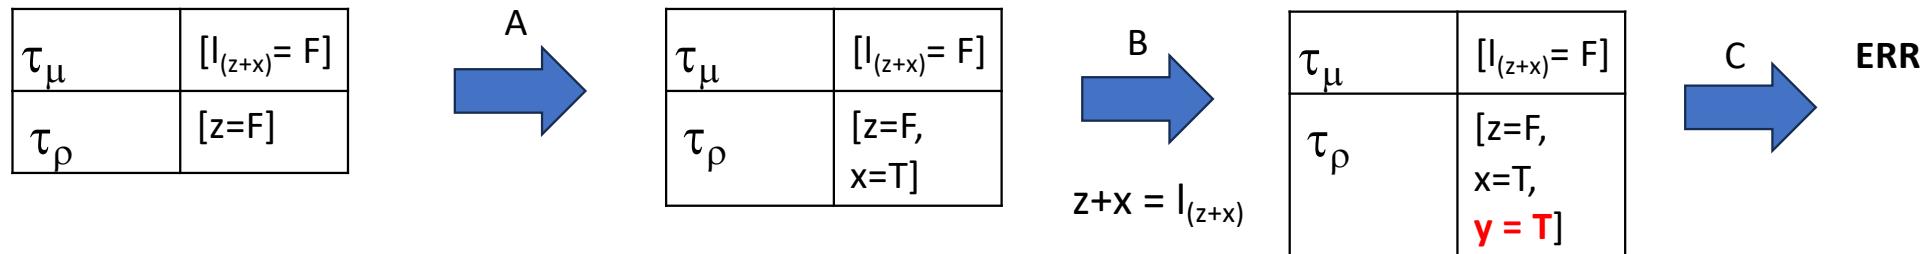
The tainted value of load is a combination of the tainted value of the location and the tainted value of the cell

Tainted jump

$$\frac{\tau_\mu, \tau_\rho, \mu, \rho \vdash e \Downarrow \langle v, t \rangle}{\tau_\mu, \tau_\rho, \mu, \rho \vdash \text{load } e \Downarrow \langle \mu(v), \text{TPmem}(t, \tau_\mu(v)) \rangle}$$

- A. `x = get_input()`
- B. `y = load(z+x)`
- C. `goto y`

$$\text{TPmem}(t_a, t_v) = t_a \text{ OR } t_v$$



$$\text{load}(x + y) \Downarrow \langle l_{\{z+x\}}, T \rangle$$

TPmem(T, $\tau_\mu(I_{(z+x)})$) TPmem(T, F) = T OR F = T



Control flow taint

```
1 x := get_input()
2 if x = 1 then goto 3 else goto 4
3 y := 1
4 z := 42
```

We make this decision in terms of a tainted data.

The assignment to y is control-dependent on line 2, since the branching outcome determines whether or not line 3 is executed.

The assignment to z is not control-dependent on line 2, since z will be assigned the value 42 regardless of which branch is taken.

There are in programs some control flows that depend on the evaluation of the condition. The choice of the 2 branches is under control of the attacker. We would need a tainted status for the program counter. This makes things much more complicated of course.

```
1  x := get_input()
2  if x = 1 then goto 3 else goto 4
3  y := 1
4  z := 42
```

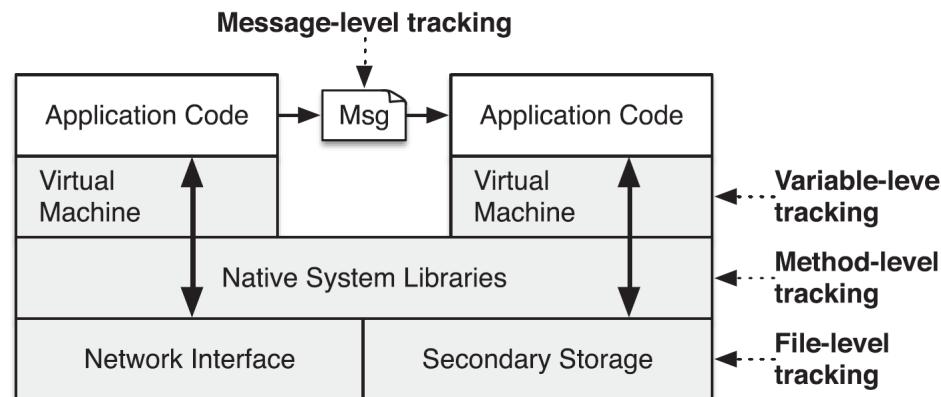
Control Flow Taint

Dynamic taint analysis cannot compute control dependencies,
thus cannot determine control-flow-based taint.

Reasoning about control dependencies requires reasoning about
multiple paths, and dynamic analysis executes on a single path at
a time.

Dynamic Taint Analysis in practice: TaintDroid

- William Enck, *et al.* “TaintDroid: An Information-Flow Tracking System for Realtime Privacy Monitoring on Smartphones,” OSDI, 2010.
- Privacy leakage (misuse) detection for Android applications.
- Employ dynamic taint analysis and report leakage during runtime.
- Challenge: Track the propagations of private data in Android platform, *e.g.*, Inter-process communication.



Taint Propagation of Interpreted Code

- Variable-level taint tracking

| Op Format | Op Semantics | Taint Propagation | Description |
|------------------------------------|----------------------------------|---------------------------------------------------------------|---------------------------------------------------|
| <i>const-op</i> $v_A \ C$ | $v_A \leftarrow C$ | $\tau(v_A) \leftarrow \emptyset$ | Clear v_A taint |
| <i>move-op</i> $v_A \ v_B$ | $v_A \leftarrow v_B$ | $\tau(v_A) \leftarrow \tau(v_B)$ | Set v_A taint to v_B taint |
| <i>move-op-R</i> v_A | $v_A \leftarrow R$ | $\tau(v_A) \leftarrow \tau(R)$ | Set v_A taint to return taint |
| <i>return-op</i> v_A | $R \leftarrow v_A$ | $\tau(R) \leftarrow \tau(v_A)$ | Set return taint (\emptyset if void) |
| <i>move-op-E</i> v_A | $v_A \leftarrow E$ | $\tau(v_A) \leftarrow \tau(E)$ | Set v_A taint to exception taint |
| <i>throw-op</i> v_A | $E \leftarrow v_A$ | $\tau(E) \leftarrow \tau(v_A)$ | Set exception taint |
| <i>unary-op</i> $v_A \ v_B$ | $v_A \leftarrow \otimes v_B$ | $\tau(v_A) \leftarrow \tau(v_B)$ | Set v_A taint to v_B taint |
| <i>binary-op</i> $v_A \ v_B \ v_C$ | $v_A \leftarrow v_B \otimes v_C$ | $\tau(v_A) \leftarrow \tau(v_B) \cup \tau(v_C)$ | Set v_A taint to v_B taint \cup v_C taint |
| <i>binary-op</i> $v_A \ v_B$ | $v_A \leftarrow v_A \otimes v_B$ | $\tau(v_A) \leftarrow \tau(v_A) \cup \tau(v_B)$ | Update v_A taint with v_B taint |
| <i>binary-op</i> $v_A \ v_B \ C$ | $v_A \leftarrow v_B \otimes C$ | $\tau(v_A) \leftarrow \tau(v_B)$ | Set v_A taint to v_B taint |
| <i>aput-op</i> $v_A \ v_B \ v_C$ | $v_B[v_C] \leftarrow v_A$ | $\tau(v_B[\cdot]) \leftarrow \tau(v_B[\cdot]) \cup \tau(v_A)$ | Update array v_B taint with v_A taint |
| <i>aget-op</i> $v_A \ v_B \ v_C$ | $v_A \leftarrow v_B[v_C]$ | $\tau(v_A) \leftarrow \tau(v_B[\cdot]) \cup \tau(v_C)$ | Set v_A taint to array and index taint |
| <i>sput-op</i> $v_A \ f_B$ | $f_B \leftarrow v_A$ | $\tau(f_B) \leftarrow \tau(v_A)$ | Set field f_B taint to v_A taint |
| <i>sget-op</i> $v_A \ f_B$ | $v_A \leftarrow f_B$ | $\tau(v_A) \leftarrow \tau(f_B)$ | Set v_A taint to field f_B taint |
| <i>iput-op</i> $v_A \ v_B \ f_C$ | $v_B(f_C) \leftarrow v_A$ | $\tau(v_B(f_C)) \leftarrow \tau(v_A)$ | Set field f_C taint to v_A taint |
| <i>iget-op</i> $v_A \ v_B \ f_C$ | $v_A \leftarrow v_B(f_C)$ | $\tau(v_A) \leftarrow \tau(v_B(f_C)) \cup \tau(v_B)$ | Set v_A taint to f_C and obj. ref. taint |

TAINT ANALYSIS: SUMMARY

- Model: *Model of programming language behaviour*
- Threat: The attacker controls some data and attempts to taint program data in order to create security vulnerabilities (buffer overflow, string attacks, injections)
- Countermeasures: Analysis to track flow of data

References

- E. J. Schwartz, T. Avgerinos, D.Brumley:
All You Ever Wanted to Know about Dynamic Taint Analysis and
Forward Symbolic Execution (but Might Have Been Afraid to Ask).IEEE
Symposium on Security and Privacy 2010: 317-331,
- The paper defines the algorithms and summarizes the critical issues
that arise when taint analyses are used in typical security context